

# DEEP Q LEARNING AND AUTOMATIC DIFFERENTIATION

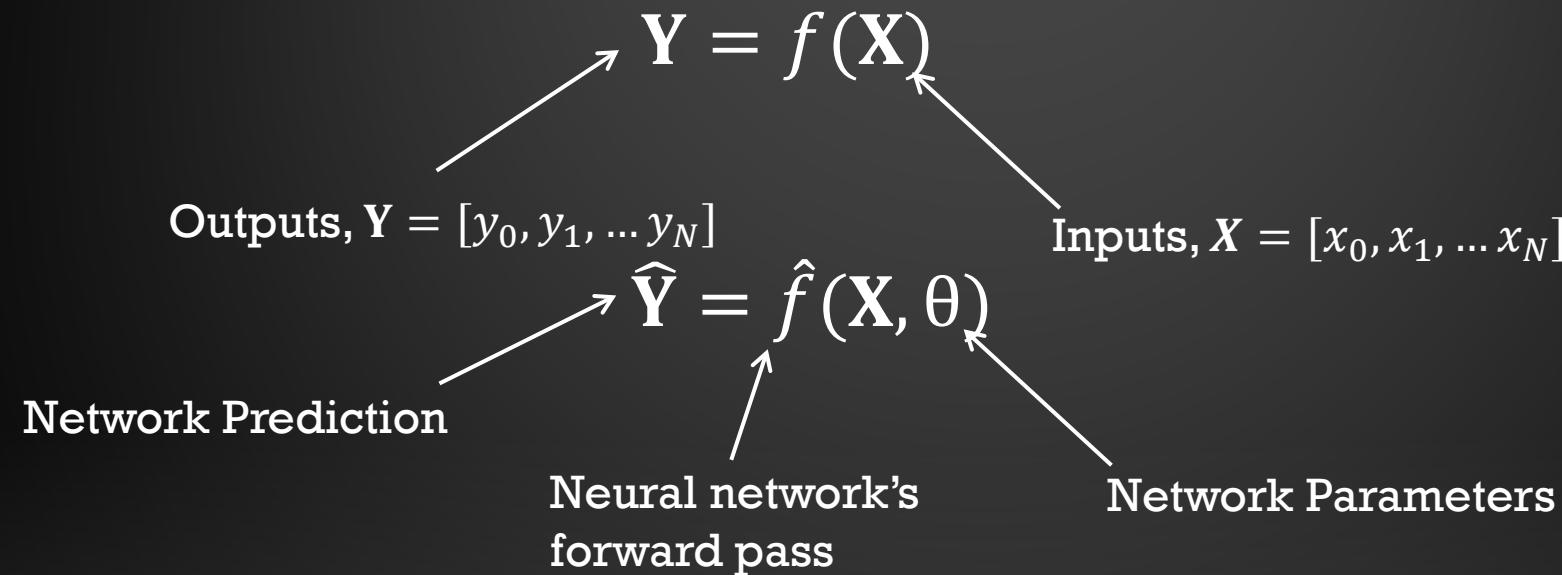
COM3240 GUEST LECTURE

IAN VIDAMOUR

# OVERVIEW

- Recap on neural networks as function approximators
- Deep Q learning
- Avoiding training instability
- Automatic differentiation
- Forward-mode AutoDiff
- Reverse-mode AutoDiff

# NEURAL NETWORKS AS FUNCTION APPROXIMATORS



# NEURAL NETWORKS AS FUNCTION APPROXIMATORS

We would like:

$$Y = \hat{Y}$$

Define a loss function:

$$\mathcal{L} \propto (Y - \hat{Y})^2$$

Change parameters  
to minimise loss:

$$\Delta\theta \propto -\frac{\partial \mathcal{L}}{\partial \theta}$$

# NEURAL NETWORKS AND REINFORCEMENT LEARNING

Q learning:

$$Q(s, a) = r(s, a) + \gamma \max_{a'} Q(s', a')$$

$$Q(s, a) = f(s, a)$$

Neural Network:

$$Q(s, a) = \hat{f}(s, a; \theta)$$

# NEURAL NETWORKS AND REINFORCEMENT LEARNING

## Update Rule:

$$Q(s, a) \leftarrow Q(s, a) + \alpha [r + \gamma \max_{a'} Q(s', a') - Q(s, a)]$$

↓  
Target    Prediction

# Predicting part of its own output?

## Nonstationary target?

# Updates biased towards single datapoints?

# MANAGING NONSTATIONARY TARGETS

Chasing a moving target causes unstable learning

Can we fix the target?

Double Q learning:

Clone the network into two networks

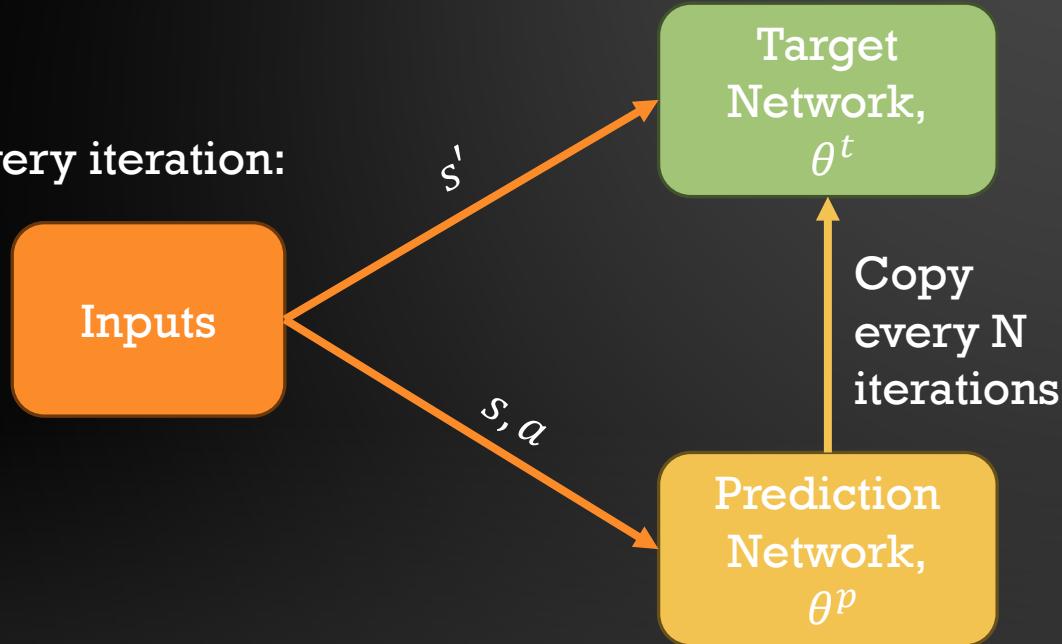
Fix one copy, use for target prediction

Update parameters of other copy to minimise loss

Update fixed network sporadically

# MANAGING NONSTATIONARY TARGETS

For every iteration:



$$\mathcal{L} = [r + \gamma \max_{a'} Q(s', a'; \theta^t) - Q(s, a; \theta^p)]^2$$

Stationary                                      Changing

# EXPLOITING PAST KNOWLEDGE

Neural networks benefit from lots of data

Batch updates smooth learning

Past information is useful!

→ Store experiences to help in future updates

# EXPLOITING PAST KNOWLEDGE

Loss function:

$$\mathcal{L} = [r + \gamma \max_{a'} Q(s', a'; \theta^t) - Q(s, a; \theta^p)]^2$$

buffer<sub>t</sub> = (s<sub>t</sub>, a<sub>t</sub>, r<sub>t</sub>, s<sub>t+1</sub>)



# AUTOMATIC DIFFERENTIATION

## Symbolic Differentiation

- + Software calculates derivatives
- + Provides general solutions
- Problems with expression swell
- Static→ Conditionals, loops, and recursion not suitable

## Numerical Differentiation

- + Conceptually simple
- Errors dependent upon step size
- Number of calculations scales poorly

# BUILDING A BUFFER

buffer<sub>t</sub> = (s<sub>t</sub>, a<sub>t</sub>, r<sub>t</sub>, s'<sub>t+1</sub>)

s	a	r	s'



# BUILDING A BUFFER

buffer<sub>t</sub> = (s<sub>t</sub>, a<sub>t</sub>, r<sub>t</sub>, s'<sub>t+1</sub>)

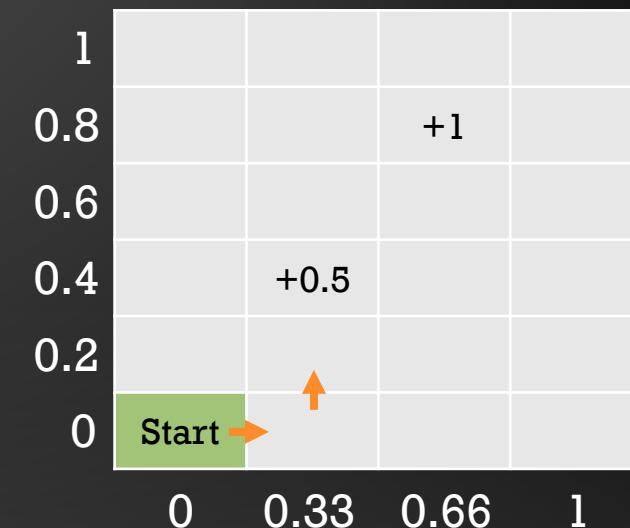
s	a	r	s'
(0,0)	(0,1,0,0)	0	(0.33,0)



# BUILDING A BUFFER

buffer<sub>t</sub> = (s<sub>t</sub>, a<sub>t</sub>, r<sub>t</sub>, s'<sub>t+1</sub>)

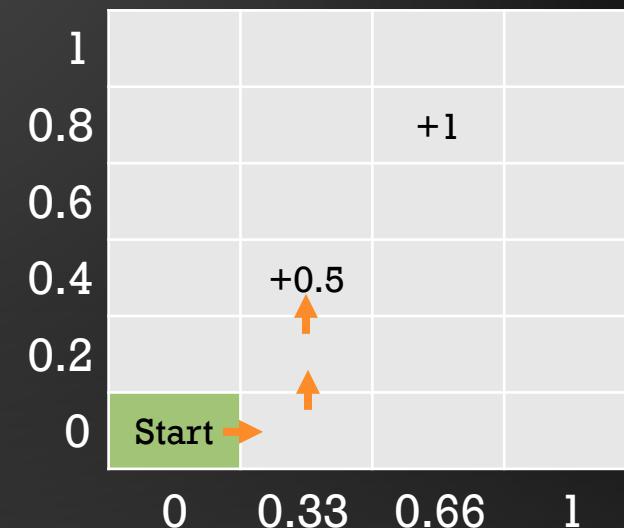
s	a	r	s'
(0,0)	(0,1,0,0)	0	(0.33,0)
(0.33,0)	(1,0,0,0)	0	(0.33,0.2)



# BUILDING A BUFFER

$\text{buffer}_t = (s_t, a_t, r_t, s_{t+1})$

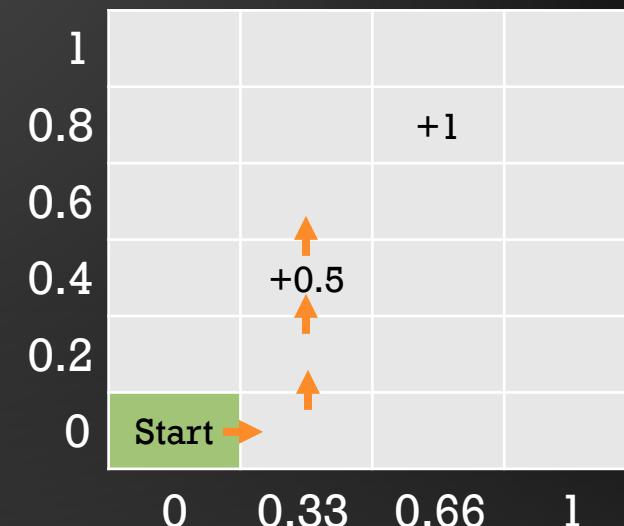
s	a	r	s'
(0,0)	(0,1,0,0)	0	(0.33,0)
(0.33,0)	(1,0,0,0)	0	(0.33,0.2)
(0.33, 0.2)	(1,0,0,0)	0.5	(0.33,0.4)



# BUILDING A BUFFER

buffer<sub>t</sub> = (s<sub>t</sub>, a<sub>t</sub>, r<sub>t</sub>, s'<sub>t+1</sub>)

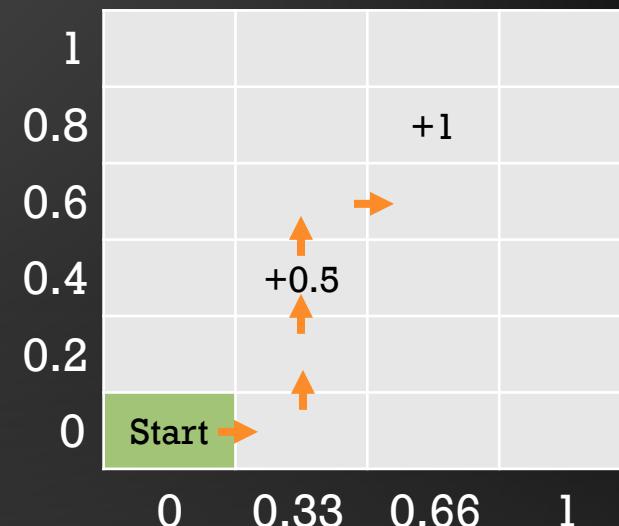
s	a	r	s'
(0,0)	(0,1,0,0)	0	(0.33,0)
(0.33,0)	(1,0,0,0)	0	(0.33,0.2)
(0.33, 0.2)	(1,0,0,0)	0.5	(0.33,0.4)
(0.33,0.4)	(1,0,0,0)	0	(0.33,0.6)



# BUILDING A BUFFER

buffer<sub>t</sub> = (s<sub>t</sub>, a<sub>t</sub>, r<sub>t</sub>, s'<sub>t+1</sub>)

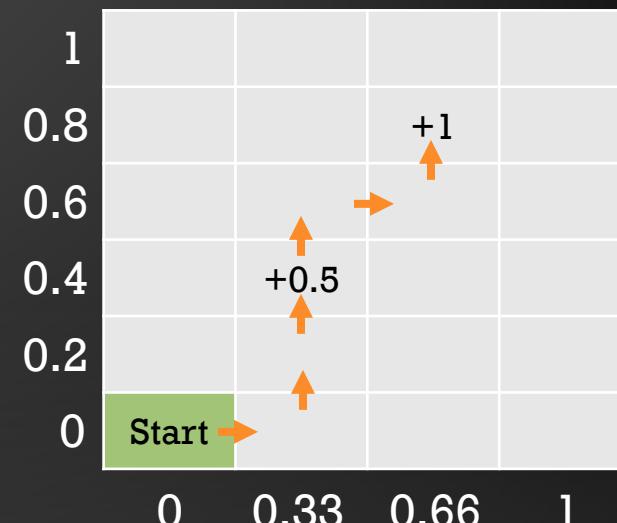
s	a	r	s'
(0,0)	(0,1,0,0)	0	(0.33,0)
(0.33,0)	(1,0,0,0)	0	(0.33,0.2)
(0.33, 0.2)	(1,0,0,0)	0.5	(0.33,0.4)
(0.33,0.4)	(1,0,0,0)	0	(0.33,0.6)
(0.33,0.6)	(0,1,0,0)	0	(0.66,0.6)



# BUILDING A BUFFER

$\text{buffer}_t = (s_t, a_t, r_t, s_{t+1})$

s	a	r	s'
(0,0)	(0,1,0,0)	0	(0.33,0)
(0.33,0)	(1,0,0,0)	0	(0.33,0.2)
(0.33, 0.2)	(1,0,0,0)	0.5	(0.33,0.4)
(0.33,0.4)	(1,0,0,0)	0	(0.33,0.6)
(0.33,0.6)	(0,1,0,0)	0	(0.66,0.6)
(0.66,0.6)	(1,0,0,0)	1	(0.66,0.8)



Continue adding in future episodes, sample for training

# EXPLOITING PAST KNOWLEDGE

Loss function:

$$\mathcal{L} = [r + \gamma \max_{a'} Q(s', a'; \theta^t) - Q(s, a; \theta^p)]^2$$

Online

Considering  
experience buffer:

$$\text{buffer}_t = (s_t, a_t, r_t, s_{t+1})$$

Batch

$$\mathcal{L} = \sum_{t=1}^N [r_t + \gamma \max_{a'} Q(s_{t+1}, a'; \theta^t) - Q(s_t, a_t; \theta^p)]^2$$

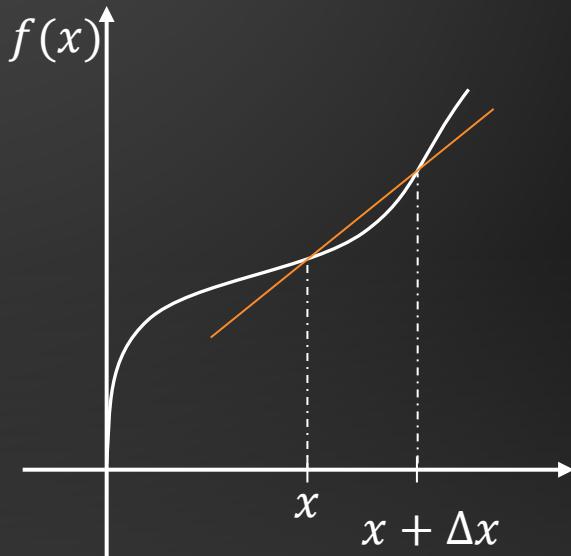
Lots of gradient calculation!

# HOW TO CALCULATE OUR DERIVATIVES

Symbolic Differentiation



Numerical Differentiation



# AUTOMATIC DIFFERENTIATION

## Automatic Differentiation

- + Pass local information numerically (avoid expression swell)
- + Resulting calculation exact
- + Synergises with backpropagation

# BUILDING COMPUTATIONAL GRAPHS

$$f(x_1, x_2) = \left[ \sin\left(\frac{x_1}{x_2}\right) + \frac{x_1}{x_2} - e^{x_2} \right] \times \left[ \frac{x_1}{x_2} - e^{x_2} \right]$$

Intermediate variables

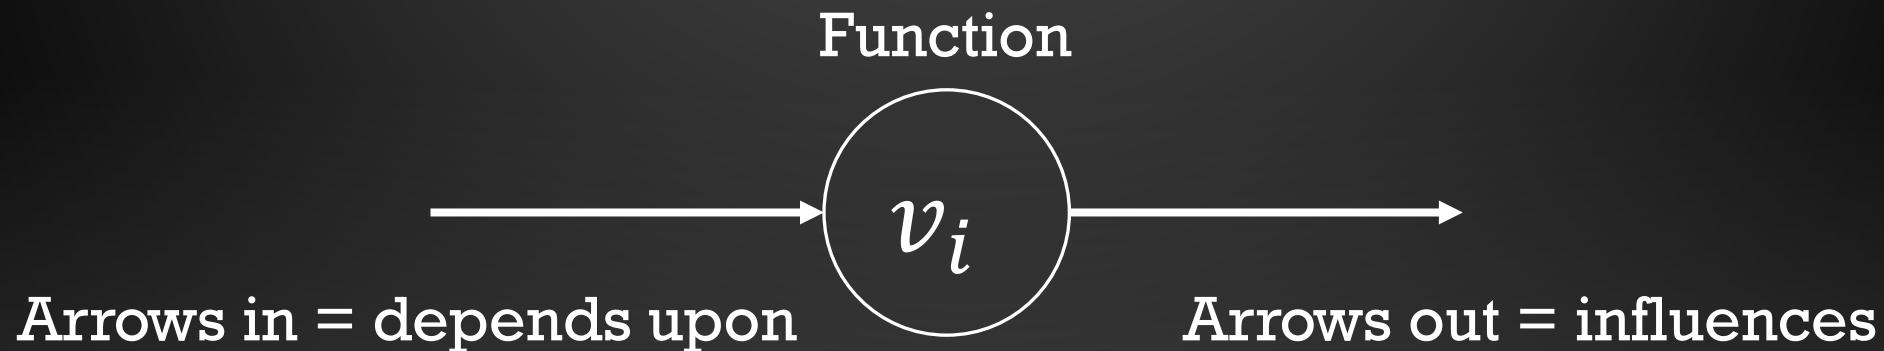
```
def f(x1, x2):
    a = (x1/x2)
    b = np.sin(a)           Primitive operations
    c = np.exp(x2)
    return (b + a - c) * (a - c)
```

# BUILDING COMPUTATIONAL GRAPHS

Computational graphs built from:

- Intermediate variables
- Primitive operations

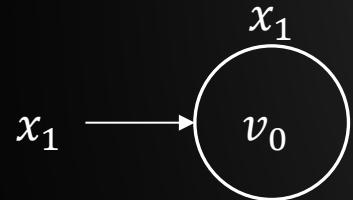
One operation per variable!



# BUILDING COMPUTATIONAL GRAPHS

$$f(x_1, x_2) = \left[ \sin\left(\frac{x_1}{x_2}\right) + \frac{x_1}{x_2} - e^{x_2} \right] \times \left[ \frac{x_1}{x_2} - e^{x_2} \right]$$

```
def f(x1, x2):
    a = (x1/x2)
    b = np.sin(a)
    c = np.exp(x2)
    return (b + a - c) * (a - c)
```

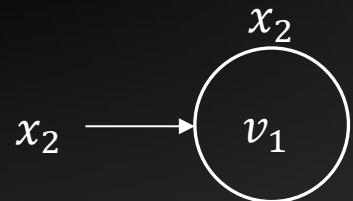
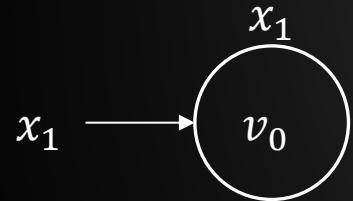


$x_2$

# BUILDING COMPUTATIONAL GRAPHS

$$f(x_1, x_2) = \left[ \sin\left(\frac{x_1}{x_2}\right) + \frac{x_1}{x_2} - e^{x_2} \right] \times \left[ \frac{x_1}{x_2} - e^{x_2} \right]$$

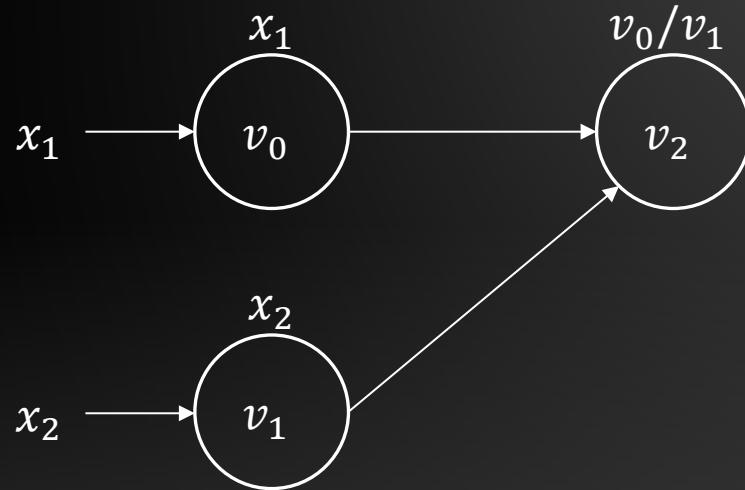
```
def f(x1, x2):
    a = (x1/x2)
    b = np.sin(a)
    c = np.exp(x2)
    return (b + a - c) * (a - c)
```



# BUILDING COMPUTATIONAL GRAPHS

$$f(x_1, x_2) = \left[ \sin\left(\frac{x_1}{x_2}\right) + \frac{x_1}{x_2} - e^{x_2} \right] \times \left[ \frac{x_1}{x_2} - e^{x_2} \right]$$

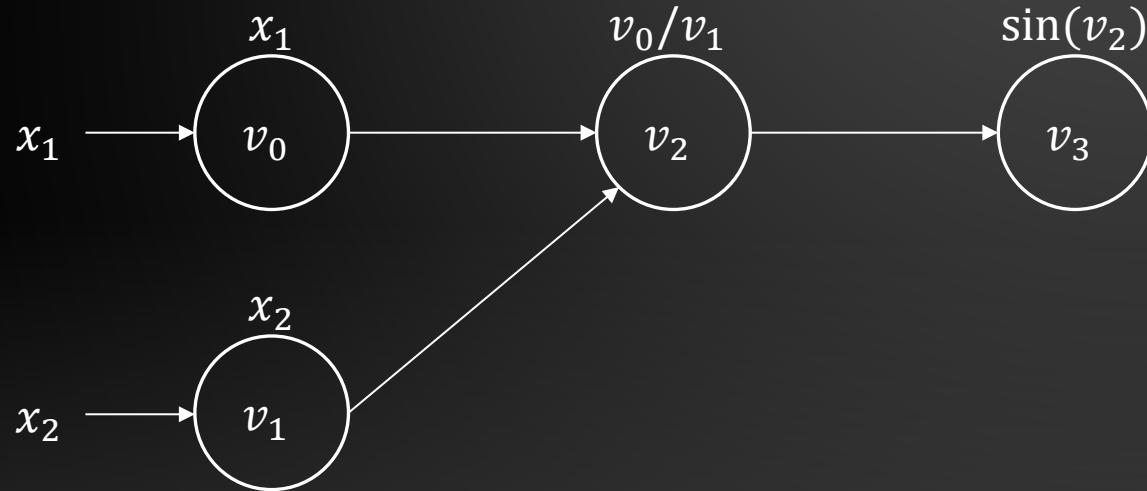
```
def f(x1, x2):
    a = (x1/x2)
    b = np.sin(a)
    c = np.exp(x2)
    return (b + a - c) * (a - c)
```



# BUILDING COMPUTATIONAL GRAPHS

$$f(x_1, x_2) = \left[ \sin\left(\frac{x_1}{x_2}\right) + \frac{x_1}{x_2} - e^{x_2} \right] \times \left[ \frac{x_1}{x_2} - e^{x_2} \right]$$

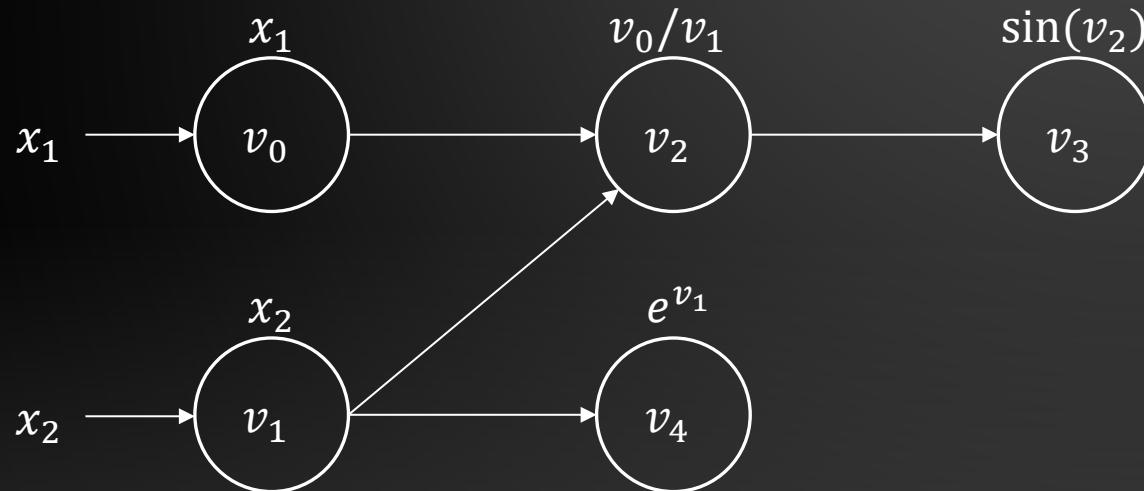
```
def f(x1, x2):
    a = (x1/x2)
    b = np.sin(a)
    c = np.exp(x2)
    return (b + a - c) * (a - c)
```



# BUILDING COMPUTATIONAL GRAPHS

$$f(x_1, x_2) = \left[ \sin\left(\frac{x_1}{x_2}\right) + \frac{x_1}{x_2} - e^{x_2} \right] \times \left[ \frac{x_1}{x_2} - e^{x_2} \right]$$

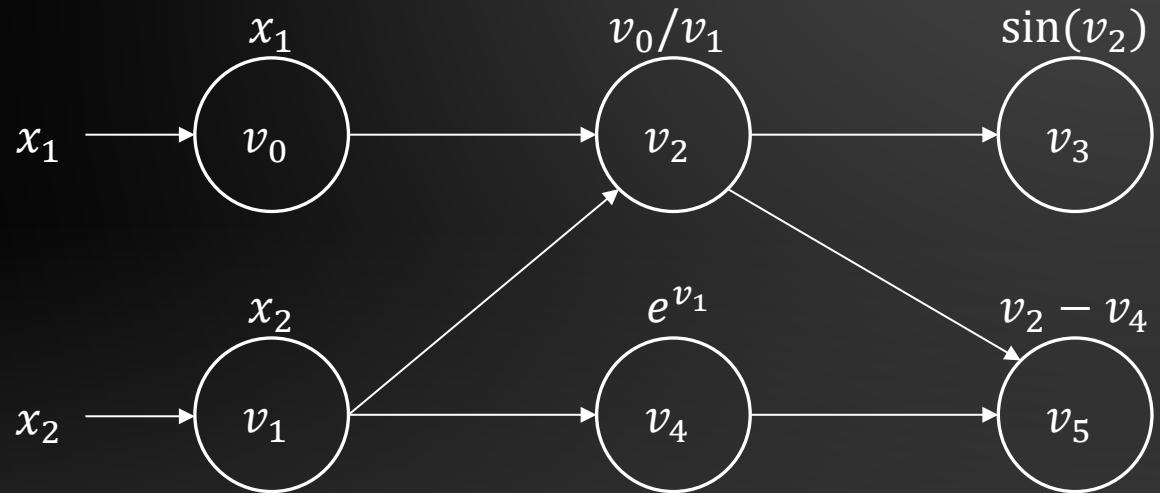
```
def f(x1, x2):
    a = (x1/x2)
    b = np.sin(a)
    c = np.exp(x2)
    return (b + a - c) * (a - c)
```



# BUILDING COMPUTATIONAL GRAPHS

$$f(x_1, x_2) = \left[ \sin\left(\frac{x_1}{x_2}\right) + \boxed{\frac{x_1}{x_2} - e^{x_2}} \right] \times \boxed{\left[ \frac{x_1}{x_2} - e^{x_2} \right]}$$

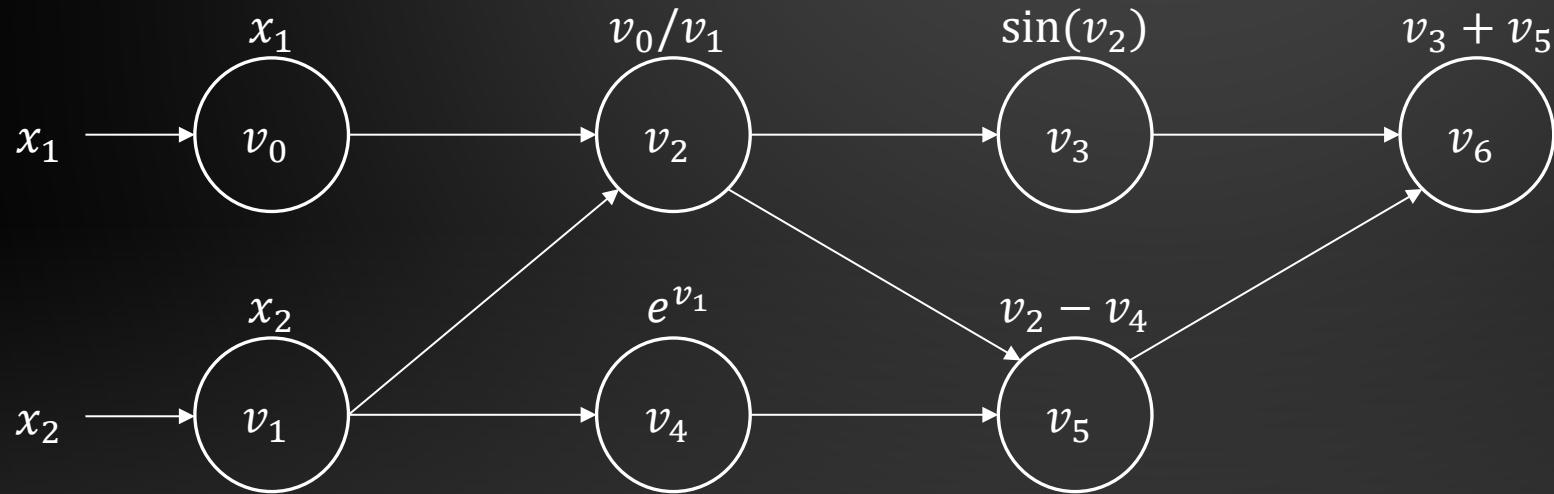
```
def f(x1, x2):
    a = (x1/x2)
    b = np.sin(a)
    c = np.exp(x2)
    return (b + a - c) * (a - c)
```



# BUILDING COMPUTATIONAL GRAPHS

$$f(x_1, x_2) = \left[ \sin\left(\frac{x_1}{x_2}\right) + \frac{x_1}{x_2} - e^{x_2} \right] \times \left[ \frac{x_1}{x_2} - e^{x_2} \right]$$

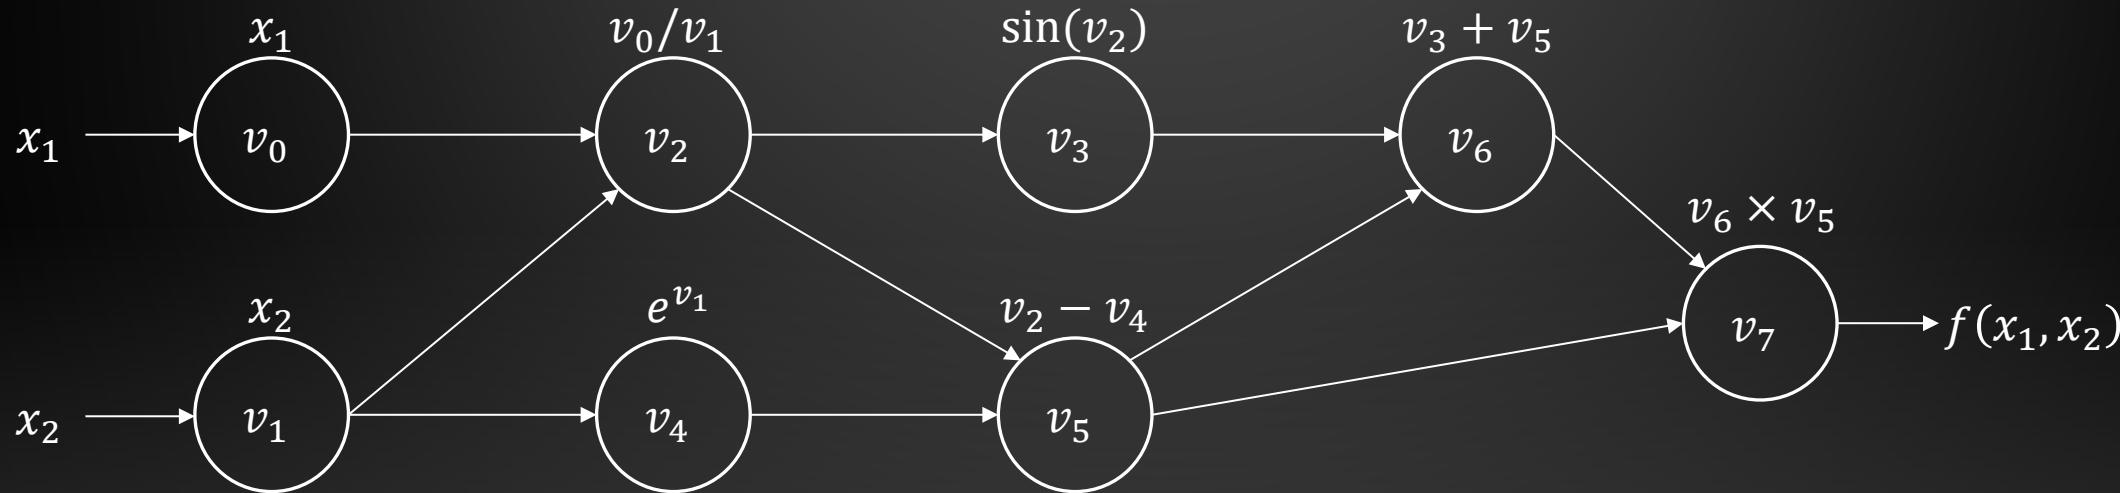
```
def f(x1, x2):
    a = (x1/x2)
    b = np.sin(a)
    c = np.exp(x2)
    return (b + a - c) * (a - c)
```



# BUILDING COMPUTATIONAL GRAPHS

$$f(x_1, x_2) = \left[ \sin\left(\frac{x_1}{x_2}\right) + \frac{x_1}{x_2} - e^{x_2} \right] \times \left[ \frac{x_1}{x_2} - e^{x_2} \right]$$

```
def f(x1, x2):
    a = (x1/x2)
    b = np.sin(a)
    c = np.exp(x2)
    return (b + a - c) * (a - c)
```



We have broken down a complex function to a series of simple operations with known derivatives

# FORWARD MODE AUTO-DIFF

Forward mode: Calculate  $x$  and  $\dot{x}$  and pass both through graph.

Perform  $\dot{x}$  operations according to gradient rules

$$\frac{\partial(\frac{u}{v})}{\partial x} = \frac{v \frac{\partial u}{\partial x} - u \frac{\partial v}{\partial x}}{v^2}$$

$$\frac{\partial(uv)}{\partial x} = v \frac{\partial u}{\partial x} + u \frac{\partial v}{\partial x}$$

$$\frac{\partial(f(g(x)))}{\partial x} = f'(g(x))g'(x)$$

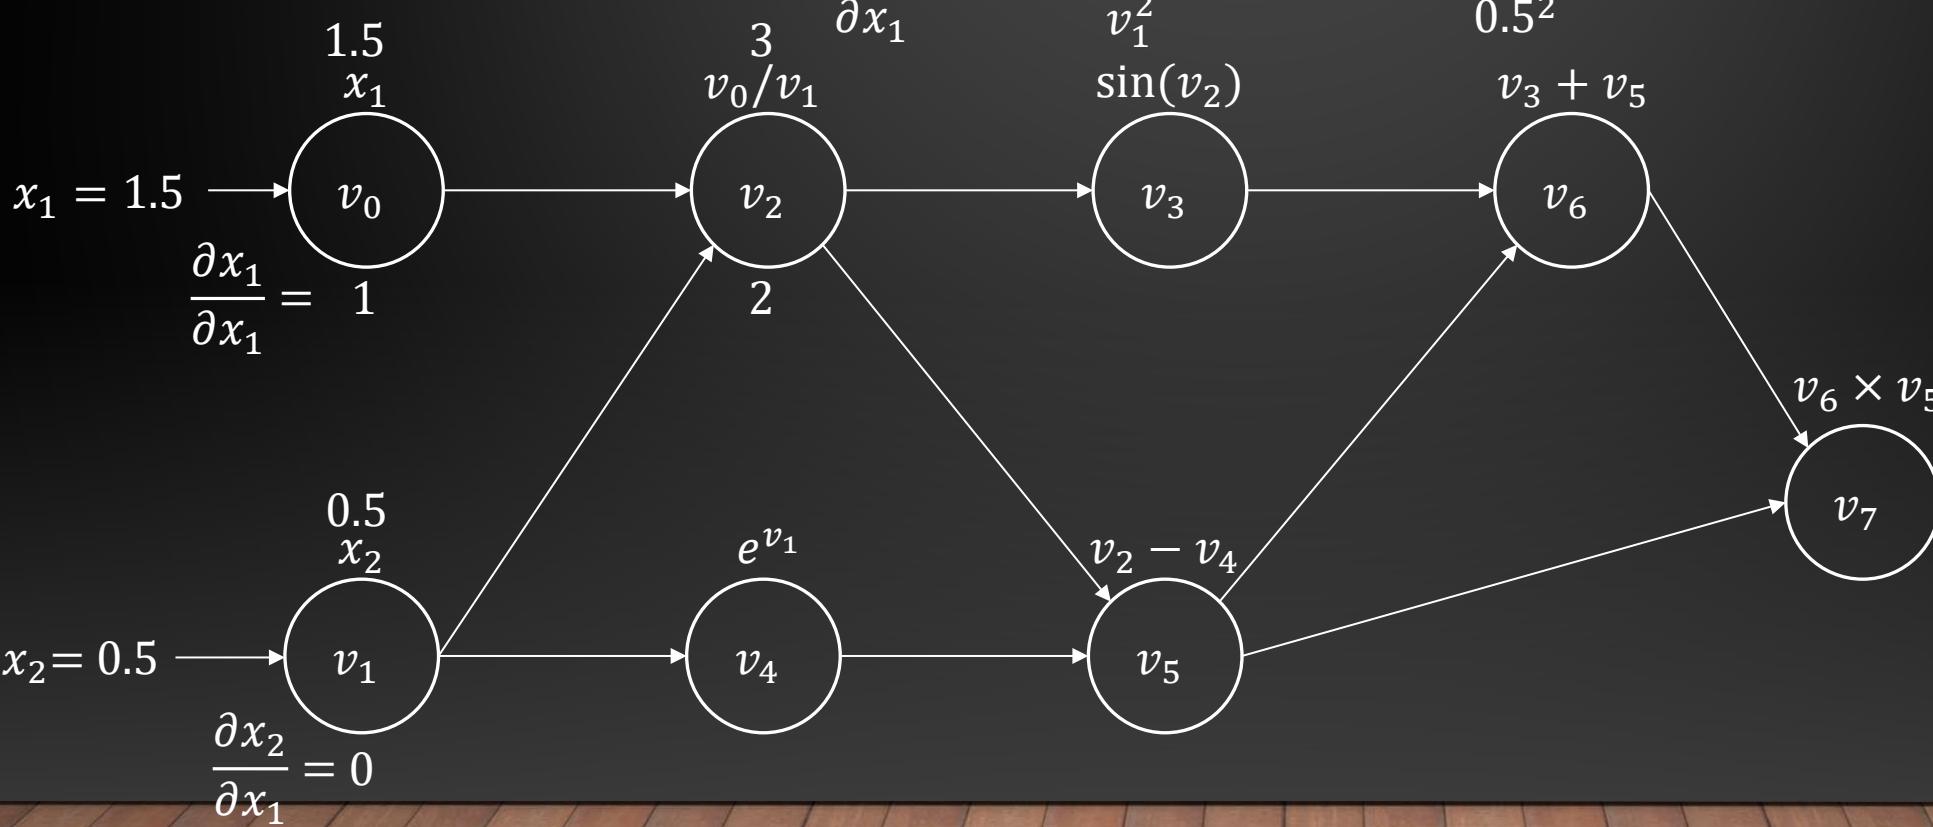
# FORWARD MODE AUTO-DIFF

$$\frac{\partial(\frac{u}{v})}{\partial x} = \frac{v \frac{\partial u}{\partial x} - u \frac{\partial v}{\partial x}}{v^2}$$

$$\frac{\partial(uv)}{\partial x} = v \frac{\partial u}{\partial x} + u \frac{\partial v}{\partial x}$$

$$\frac{\partial(f(g(x)))}{\partial x} = f'(g(x))g'(x)$$

Considering  $\frac{\partial f}{\partial x_1}$ :



function value  
 $v$   
gradient value  
 $v'$

# FORWARD MODE AUTO-DIFF

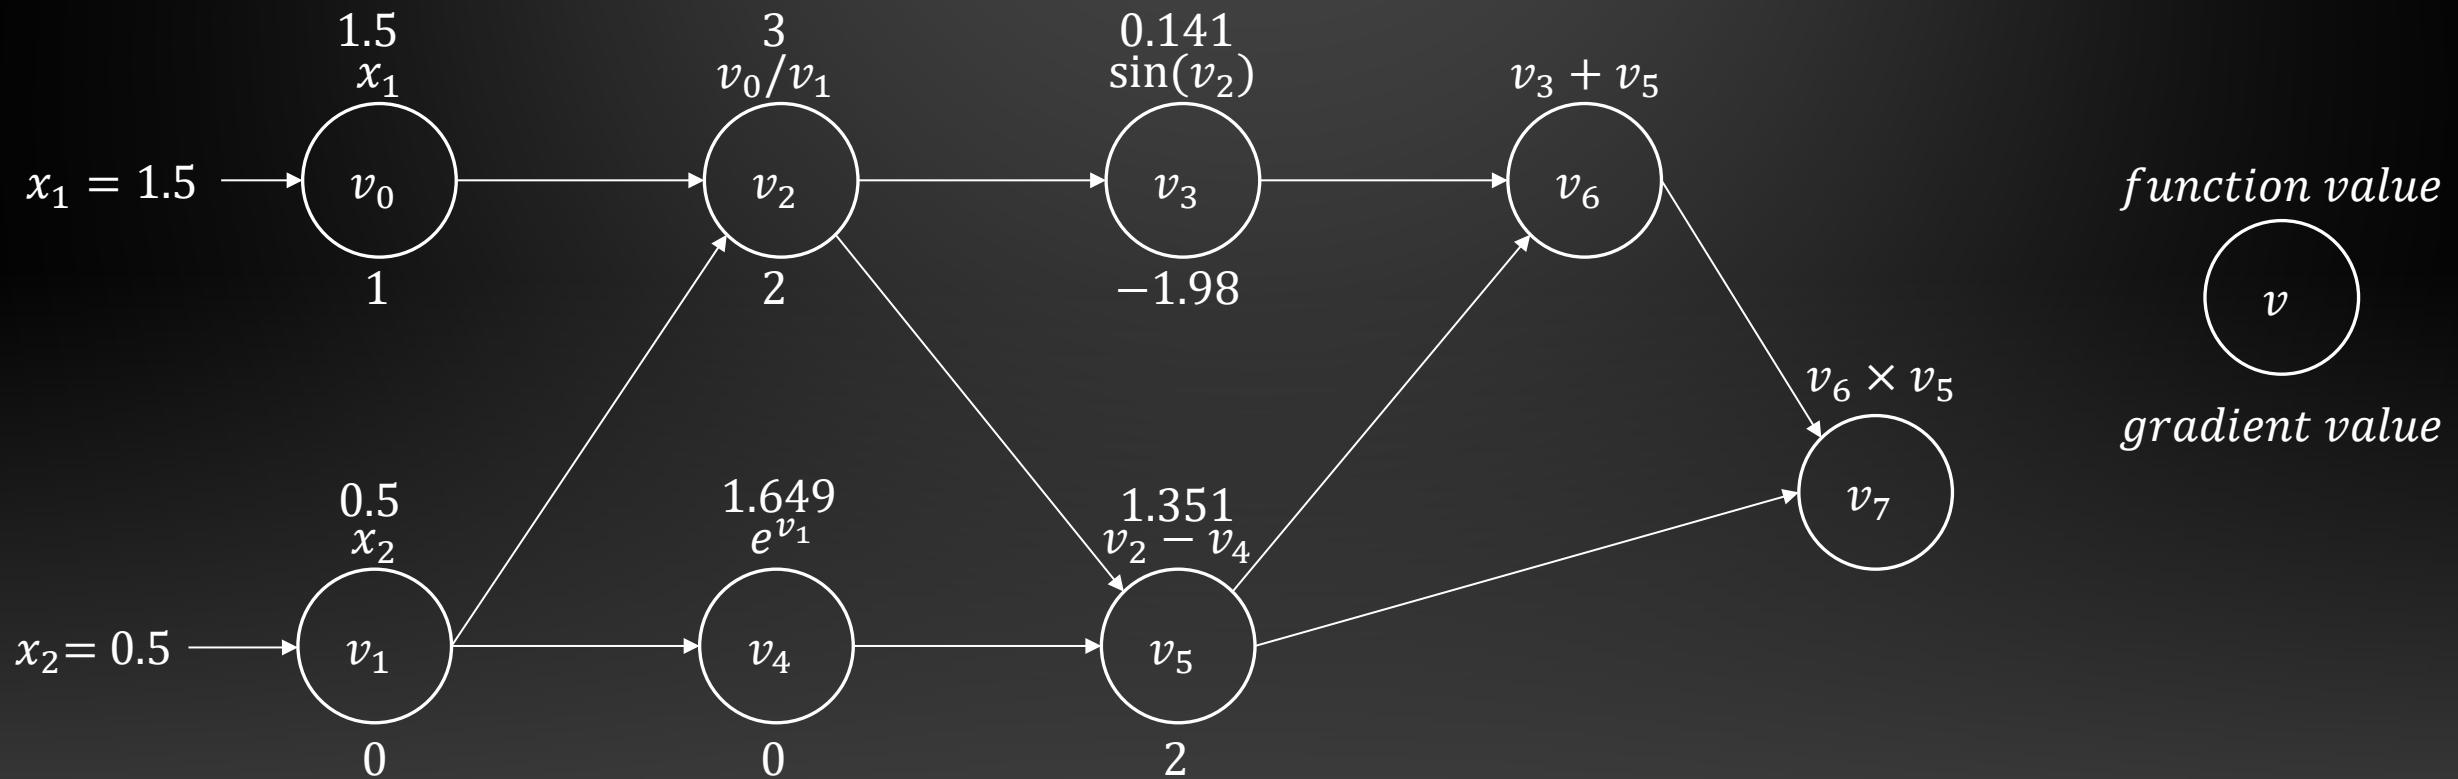
$$\frac{\partial(\frac{u}{v})}{\partial x} = \frac{v \frac{\partial u}{\partial x} - u \frac{\partial v}{\partial x}}{v^2}$$

$$\frac{\partial(uv)}{\partial x} = v \frac{\partial u}{\partial x} + u \frac{\partial v}{\partial x}$$

$$\frac{\partial(f(g(x)))}{\partial x} = f'(g(x))g'(x)$$

Considering  $\frac{\partial f}{\partial x_1}$ :

$$\frac{\partial v_3}{\partial x_1} = \frac{\partial v_2}{\partial x_1} * \cos(v_2) = 2 * \cos(3) = -1.98$$



$$\frac{\partial v_4}{\partial x_1} = 0$$

$$\frac{\partial v_5}{\partial x_1} = \frac{\partial v_2}{\partial x_1} = 2$$

# FORWARD MODE AUTO-DIFF

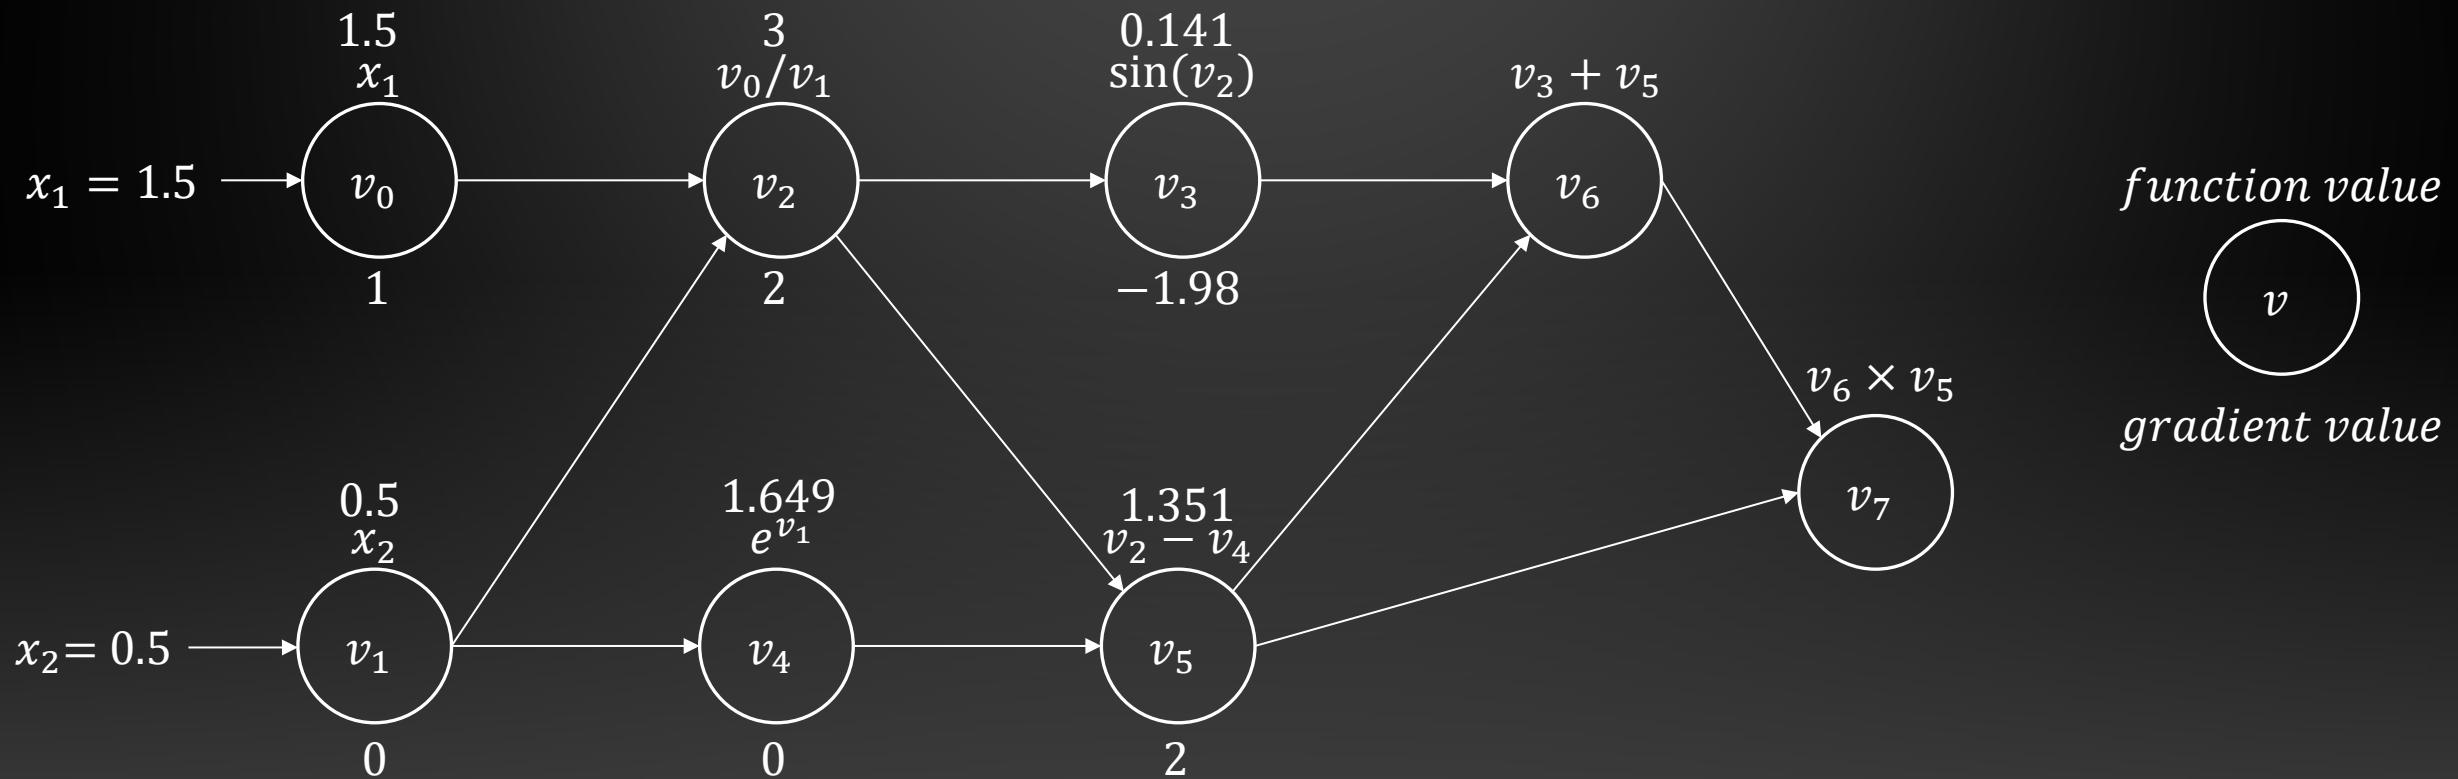
$$\frac{\partial(\frac{u}{v})}{\partial x} = \frac{v \frac{\partial u}{\partial x} - u \frac{\partial v}{\partial x}}{v^2}$$

$$\frac{\partial(uv)}{\partial x} = v \frac{\partial u}{\partial x} + u \frac{\partial v}{\partial x}$$

$$\frac{\partial(f(g(x)))}{\partial x} = f'(g(x))g'(x)$$

Considering  $\frac{\partial f}{\partial x_1}$ :

$$\frac{\partial v_3}{\partial x_1} = \frac{\partial v_2}{\partial x_1} * \cos(v_2) = 2 * \cos(3) = -1.98$$



$$\frac{\partial v_4}{\partial x_1} = 0$$

$$\frac{\partial v_5}{\partial x_1} = \frac{\partial v_2}{\partial x_1} = 2$$

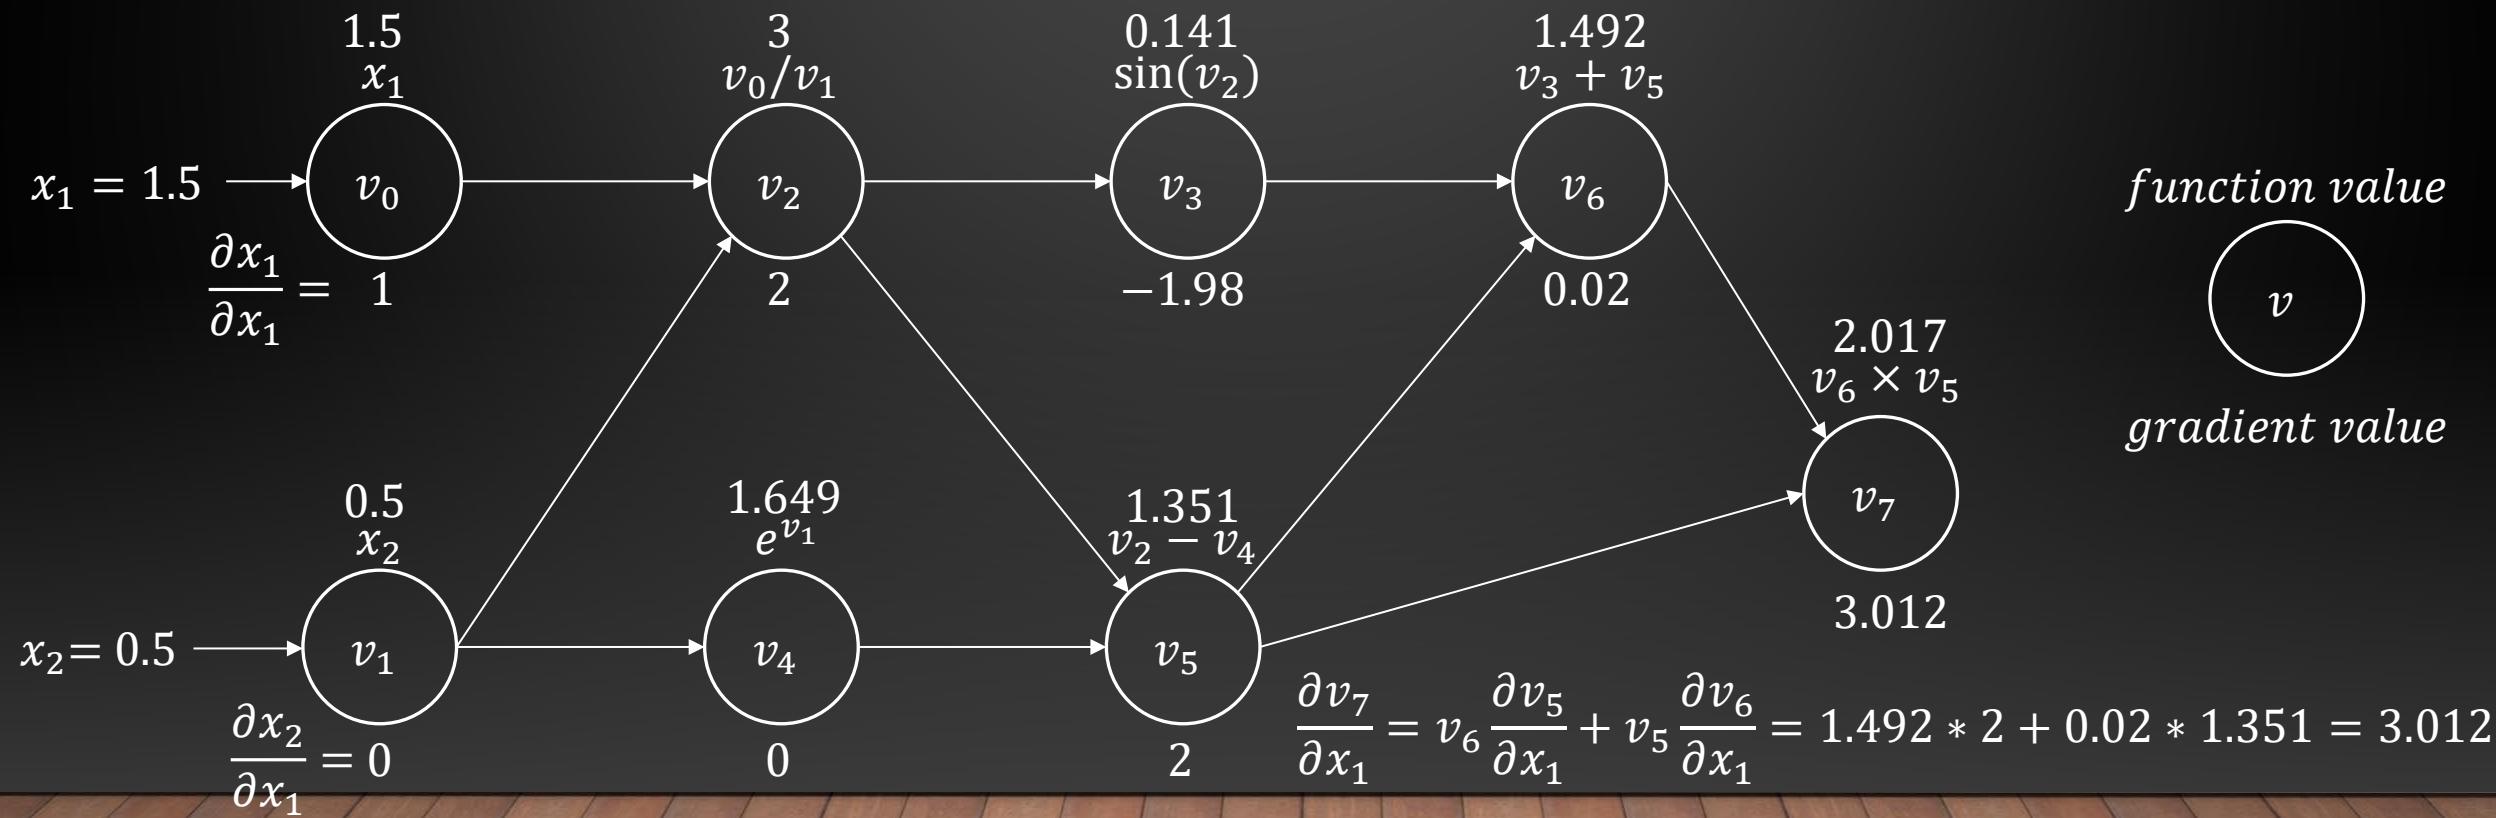
# FORWARD MODE AUTO-DIFF

$$\frac{\partial(\frac{u}{v})}{\partial x} = \frac{v \frac{\partial u}{\partial x} - u \frac{\partial v}{\partial x}}{v^2}$$

$$\frac{\partial(uv)}{\partial x} = v \frac{\partial u}{\partial x} + u \frac{\partial v}{\partial x}$$

$$\frac{\partial(f(g(x)))}{\partial x} = f'(g(x))g'(x)$$

Considering  $\frac{\partial f}{\partial x_1}$ :



# FORWARD MODE AUTO-DIFF

Consider one input at a time

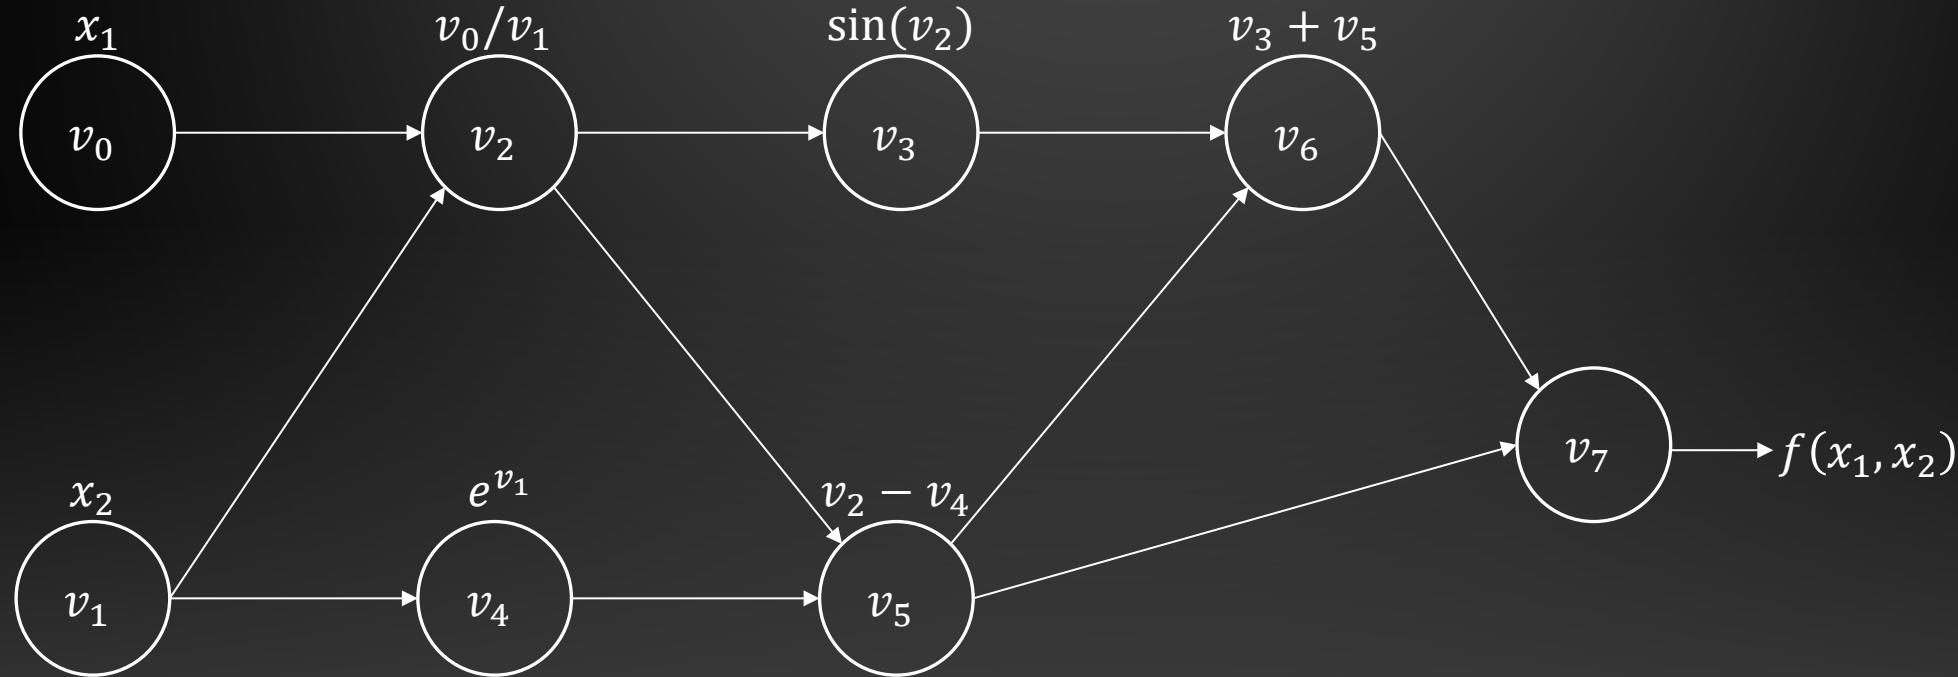
Considering many outputs from one input simple

Useful when number of N outputs > N inputs

# REVERSE MODE AUTO-DIFF

Reverse mode: pass  $x$  normally, and store dependencies (computational graph) in memory.  
Propagate derivatives backwards from the output. (Very similar to backpropagation!)

At each point in the network, we consider how all pathways ahead depend upon each intermediate variable:



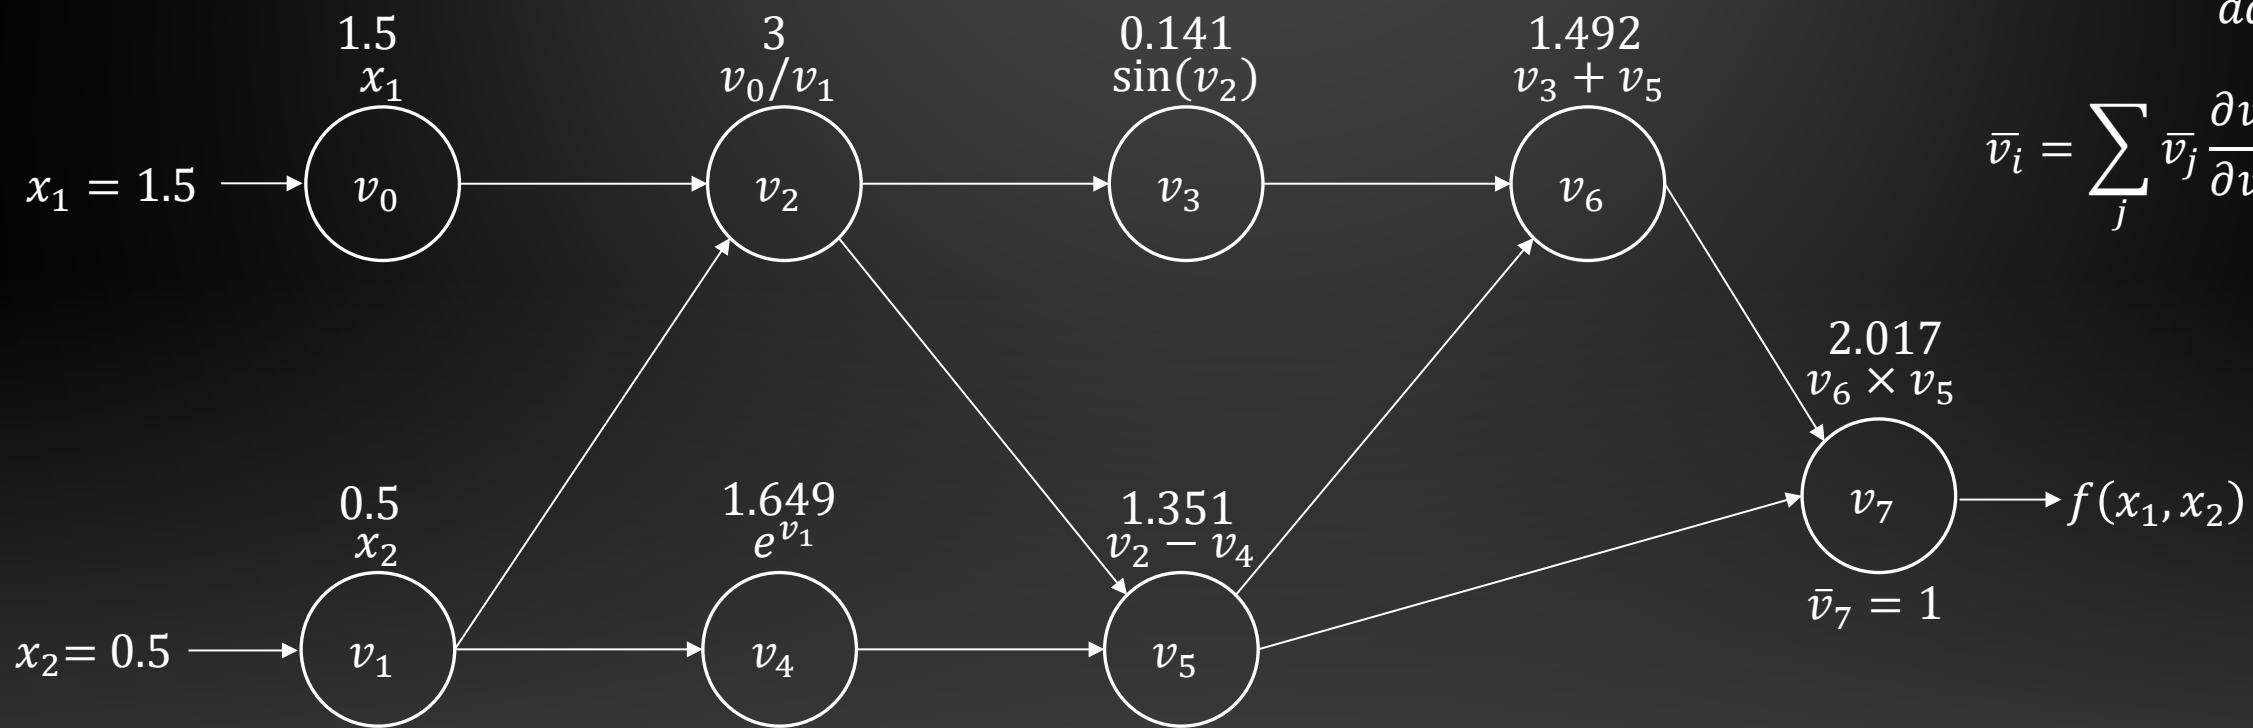
# REVERSE MODE AUTO-DIFF

function value

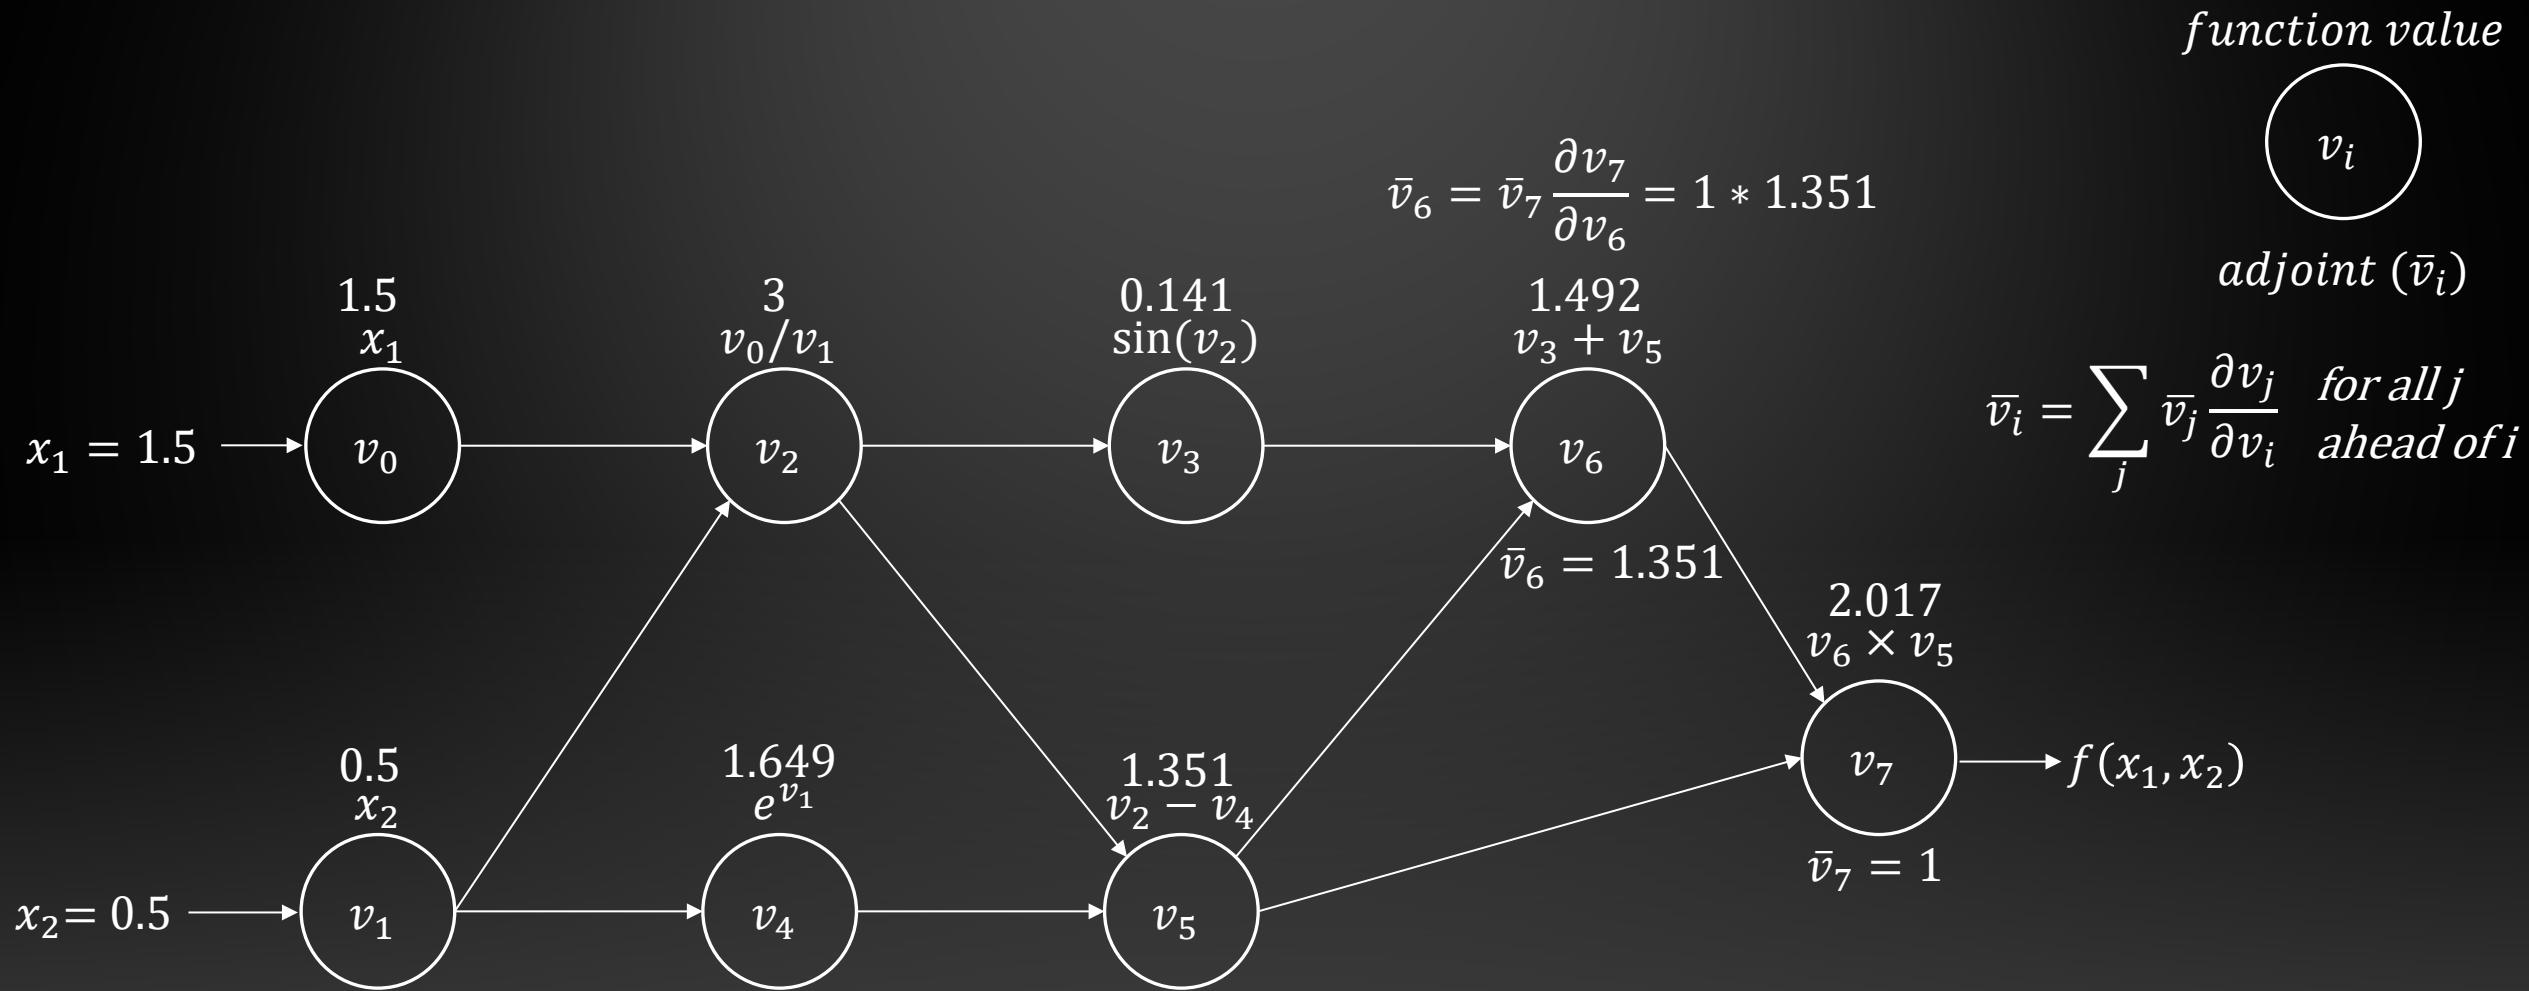


adjoint ( $\bar{v}_i$ )

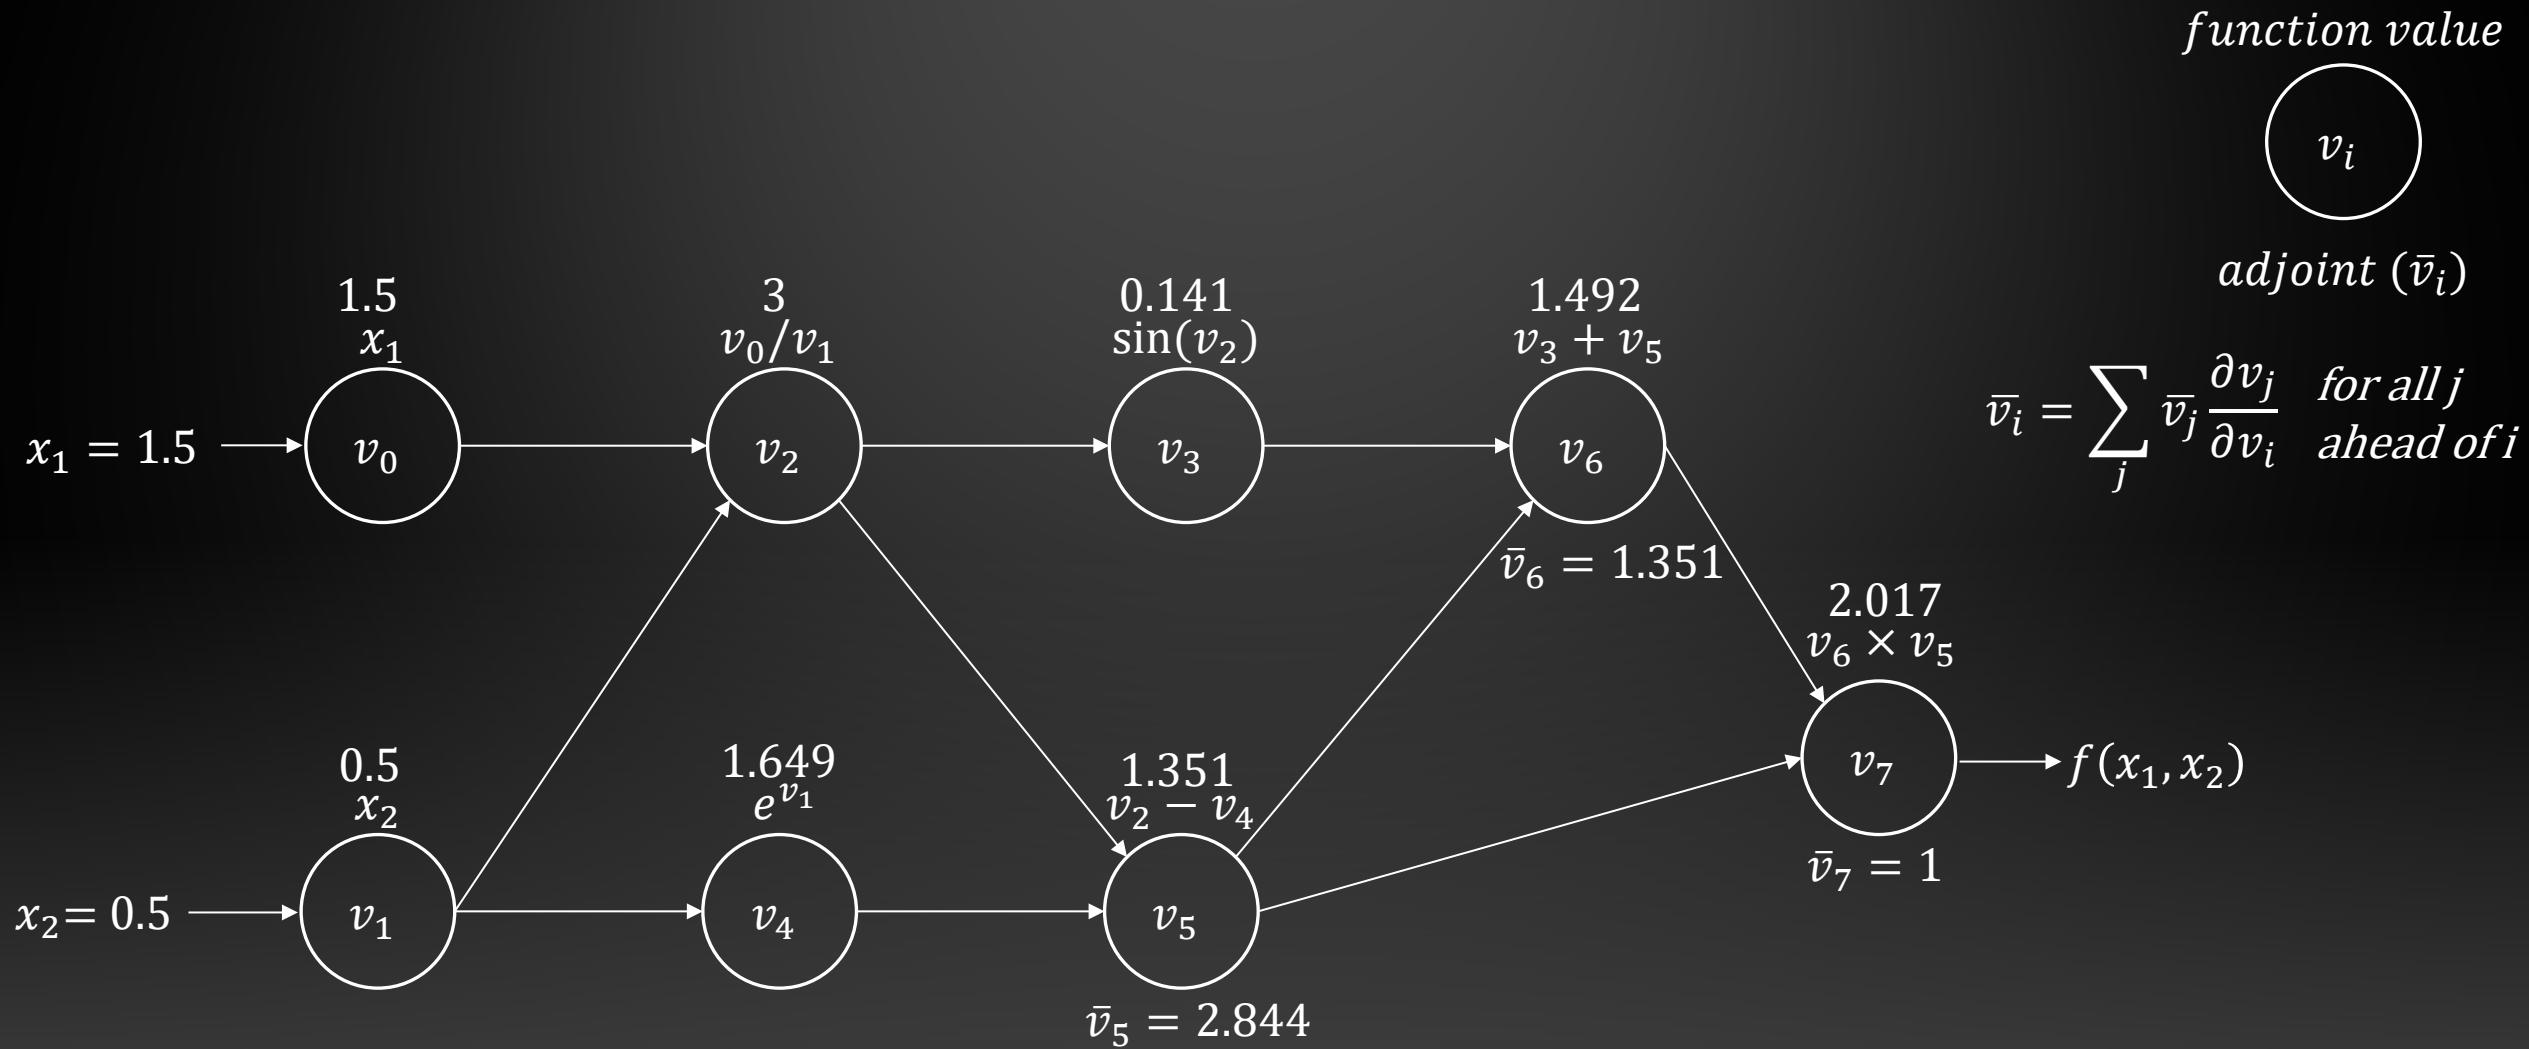
$$\bar{v}_i = \sum_j \bar{v}_j \frac{\partial v_j}{\partial v_i} \quad \text{for all } j \text{ ahead of } i$$



# REVERSE MODE AUTO-DIFF

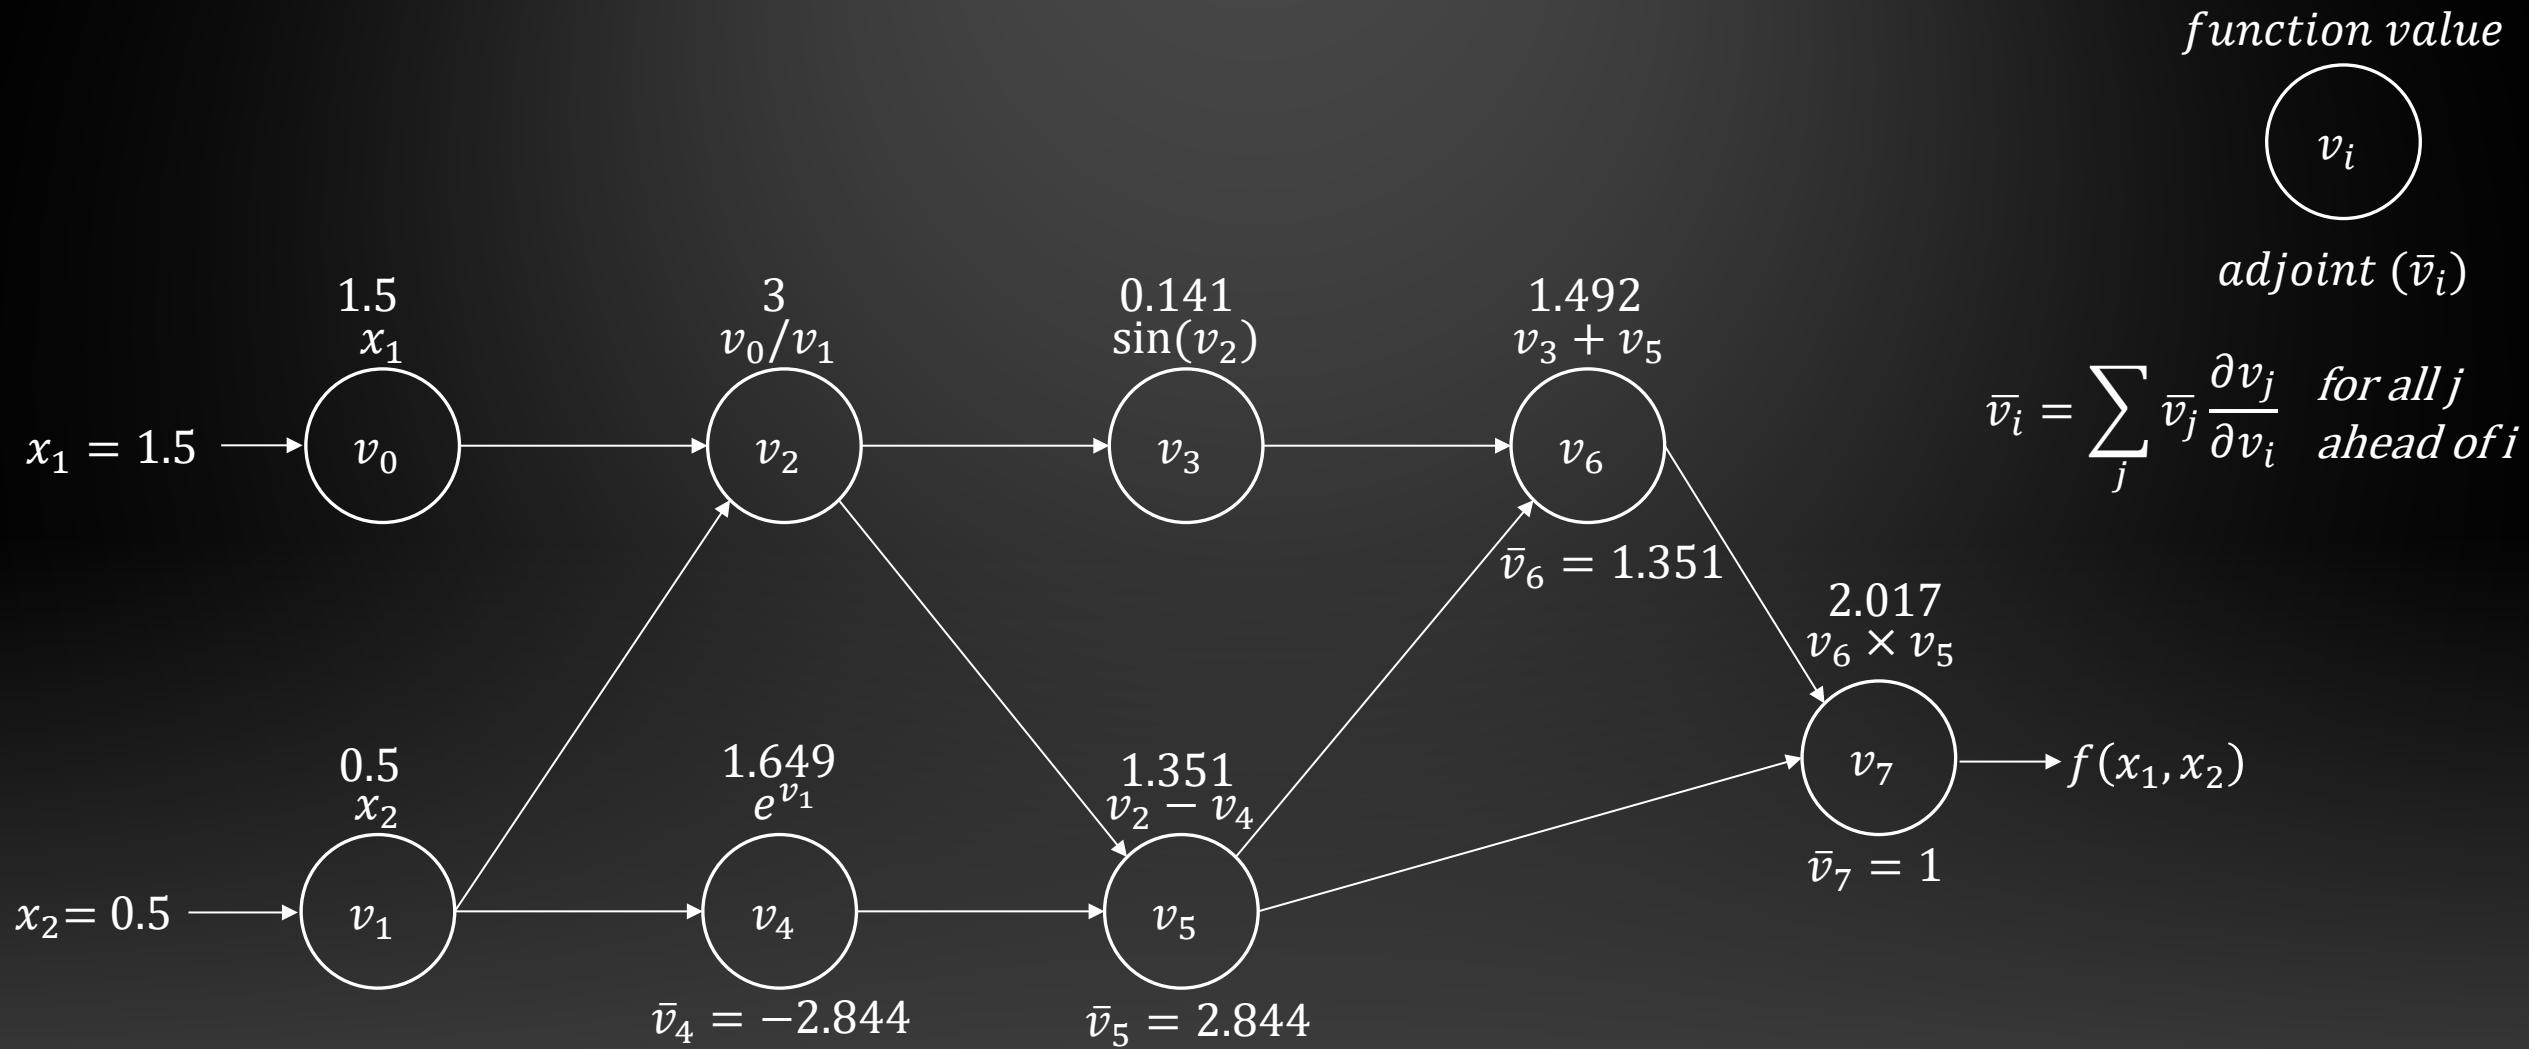


# REVERSE MODE AUTO-DIFF



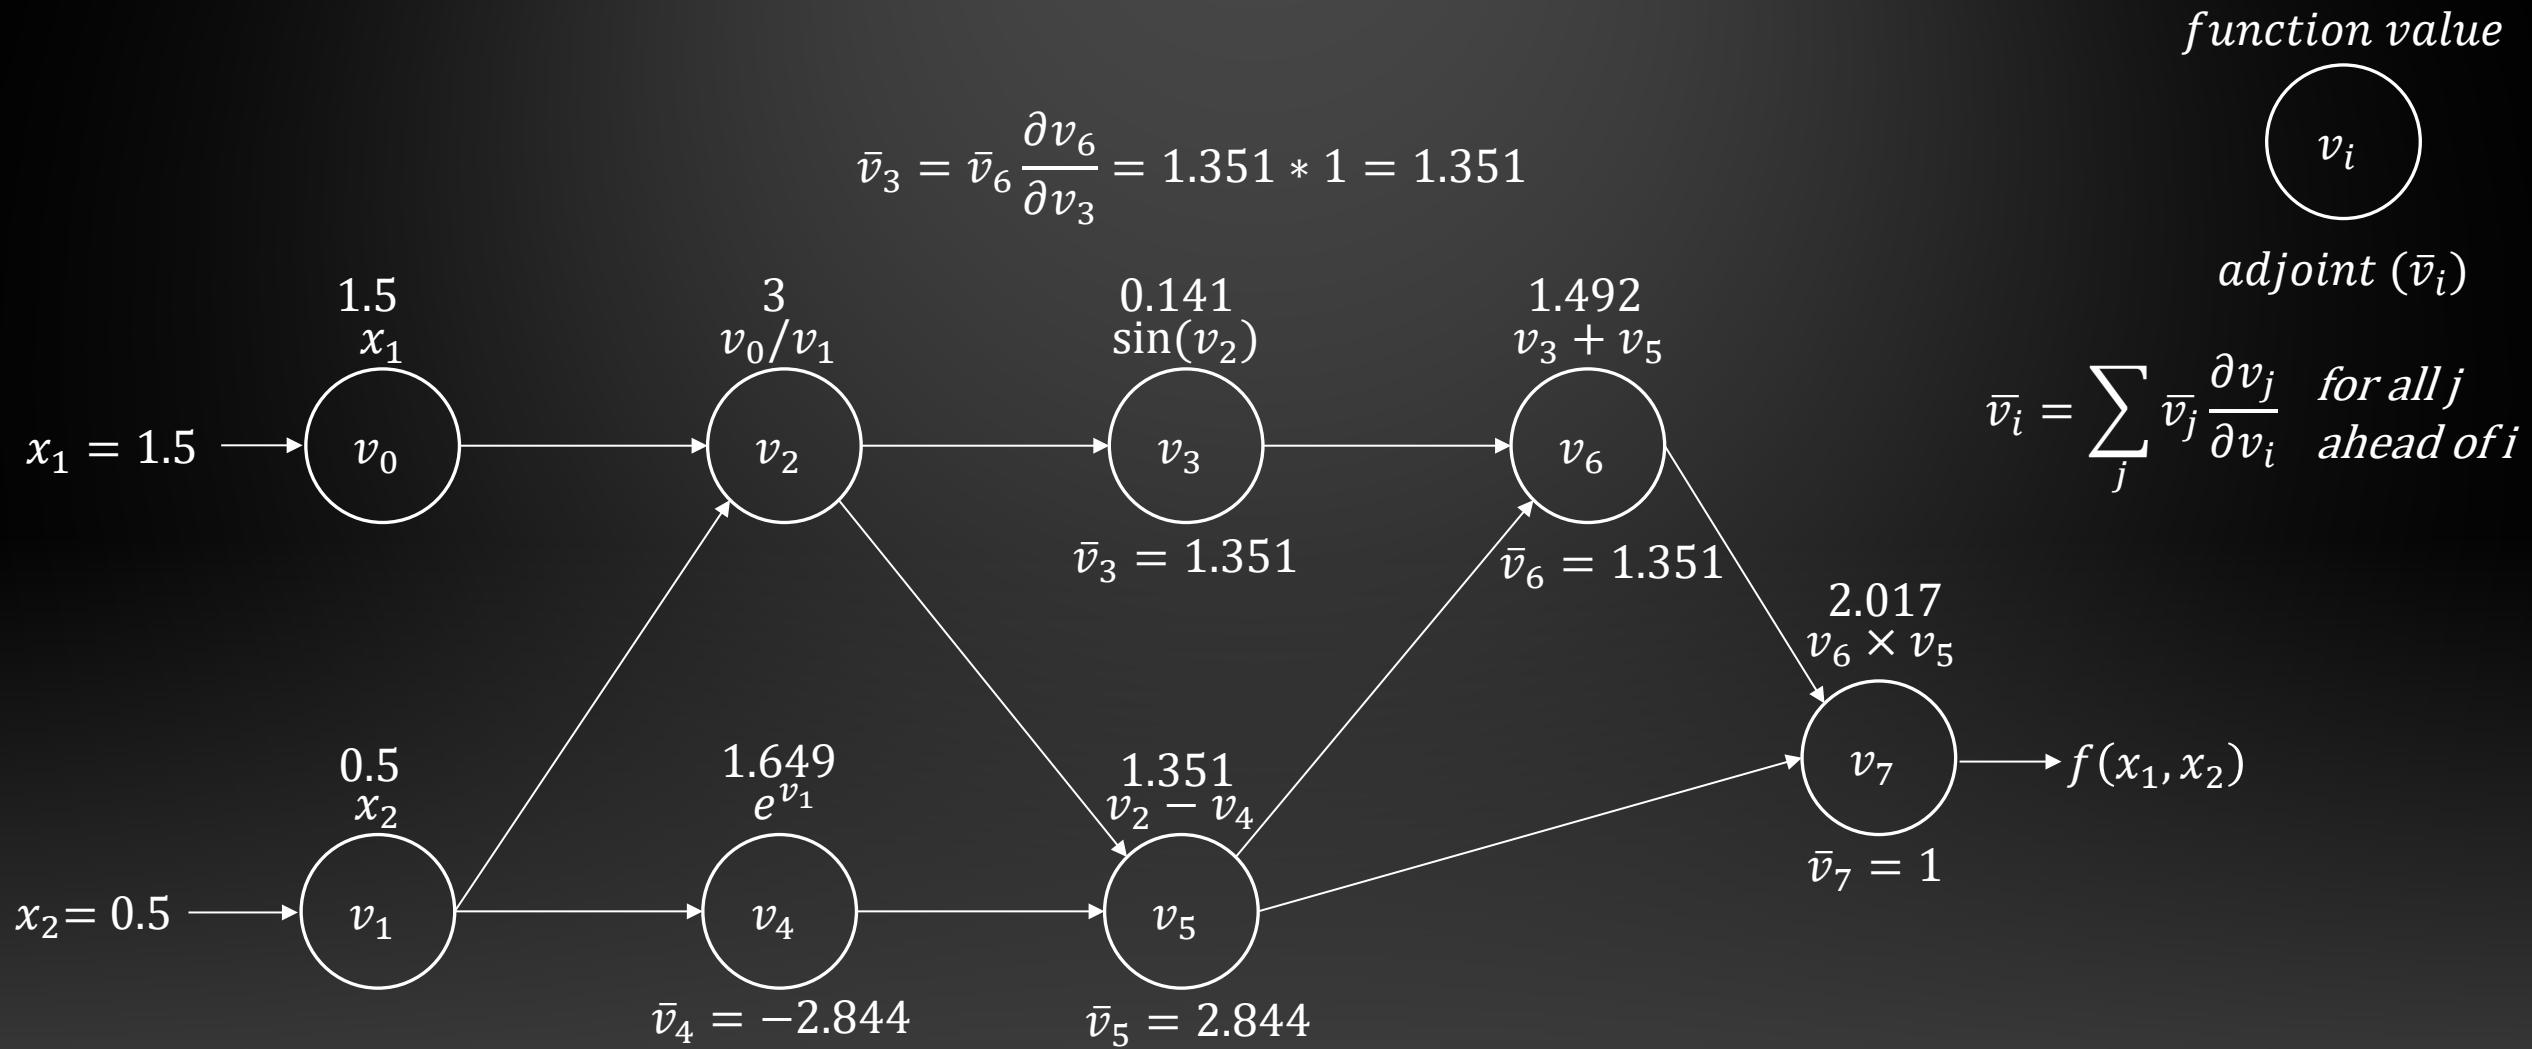
$$\bar{v}_5 = \bar{v}_7 \frac{\partial v_7}{\partial v_5} + \bar{v}_6 \frac{\partial v_6}{\partial v_5} = 1 * 1.492 + 1.351 * 1 = 2.844$$

# REVERSE MODE AUTO-DIFF

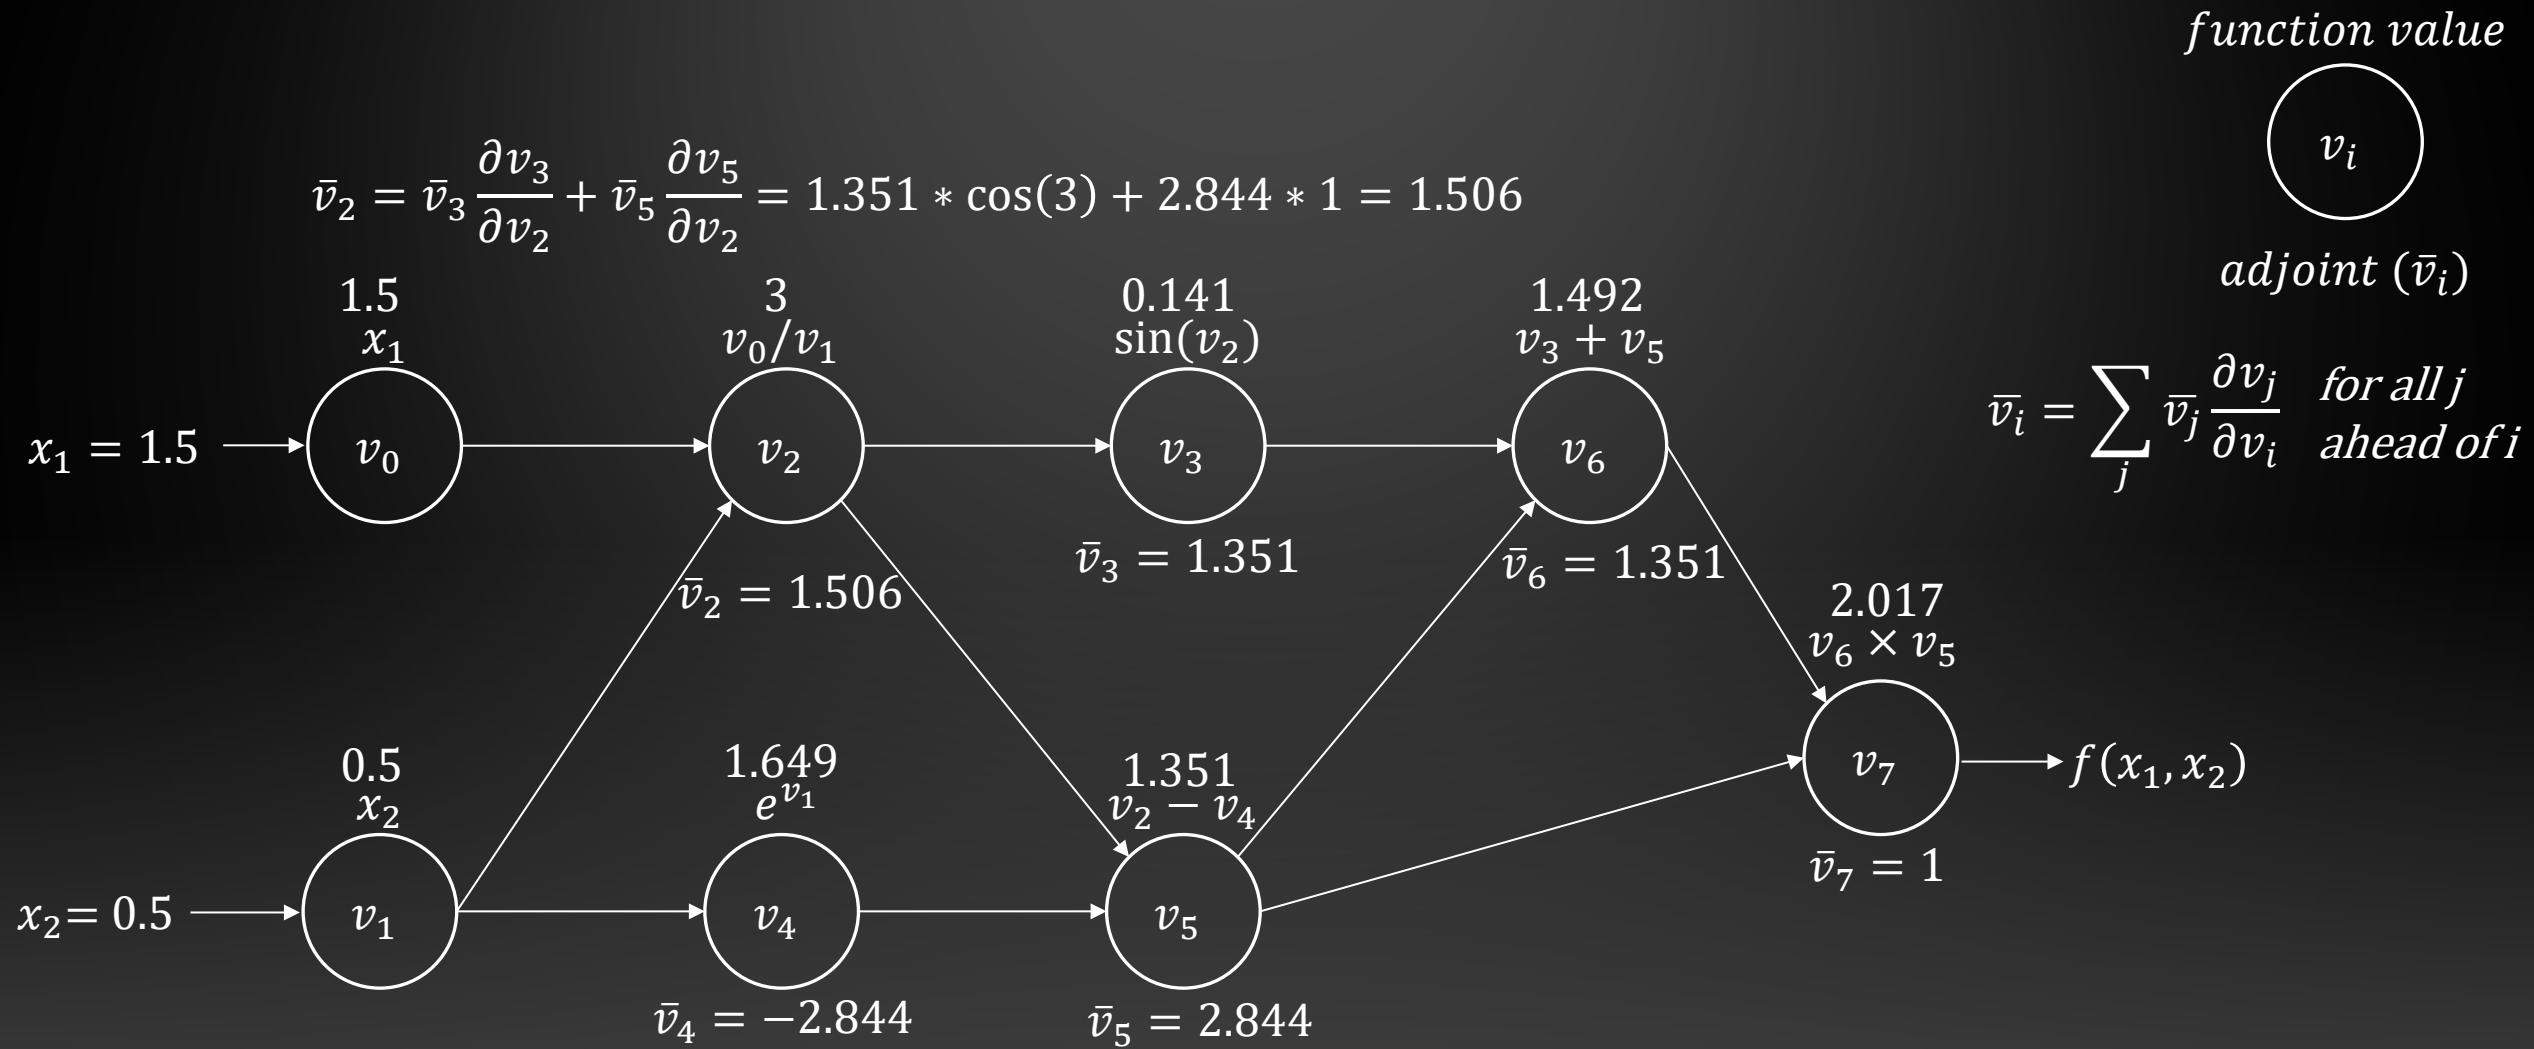


$$\bar{v}_4 = \bar{v}_5 \frac{\partial v_5}{\partial v_4} = 2.844 * -1$$

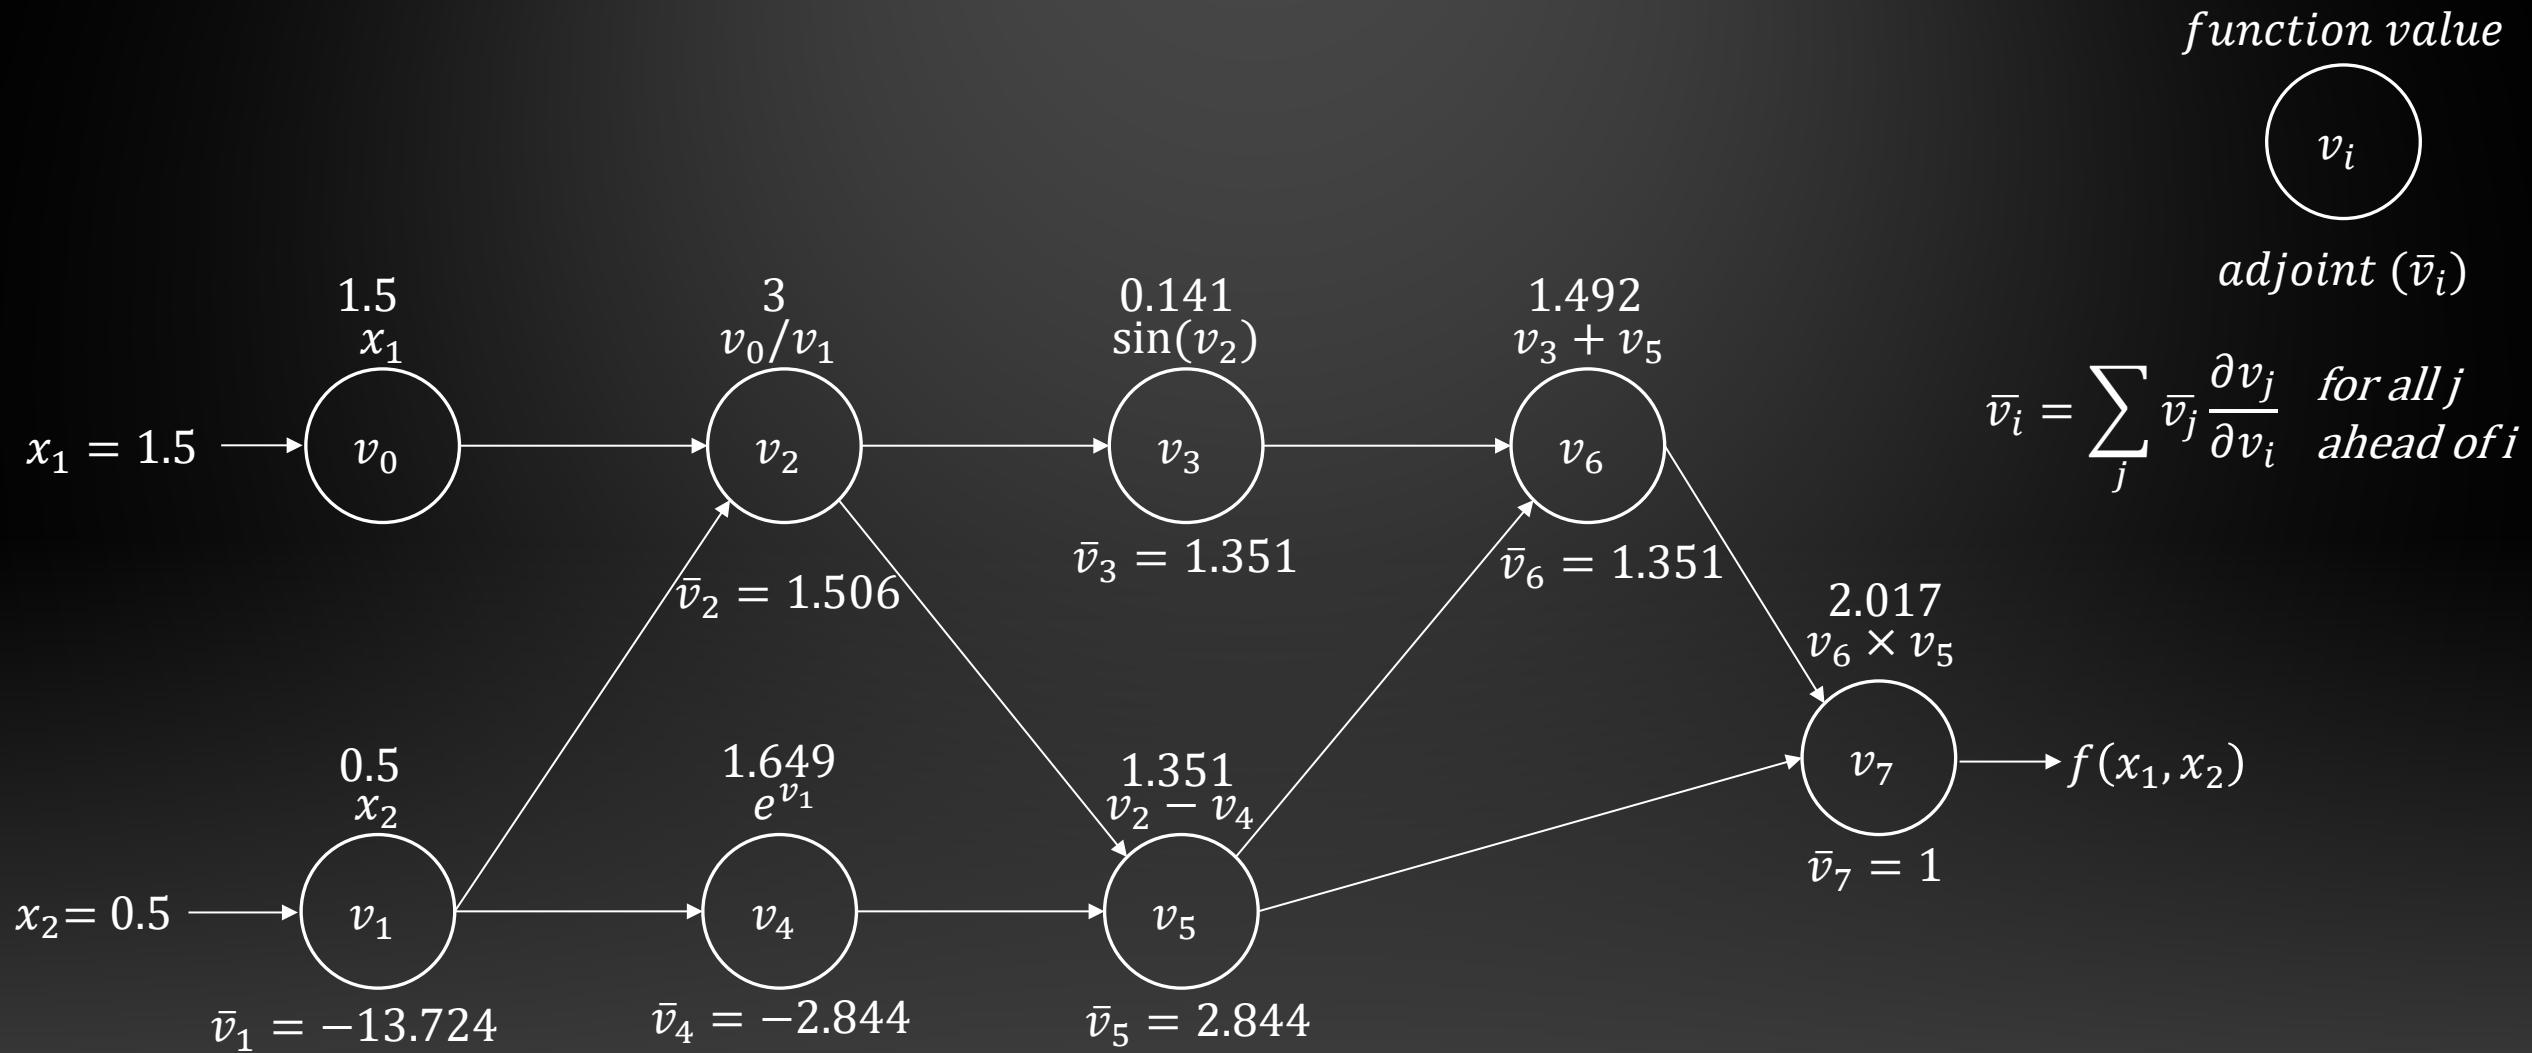
# REVERSE MODE AUTO-DIFF



# REVERSE MODE AUTO-DIFF



# REVERSE MODE AUTO-DIFF



$$\bar{v}_1 = \bar{v}_2 \frac{\partial v_2}{\partial v_1} + \bar{v}_4 \frac{\partial v_4}{\partial v_1} = 1.506 * -\frac{1.5}{0.5^2} - 2.844 * e^{0.5} = -13.724$$

# REVERSE MODE AUTO-DIFF

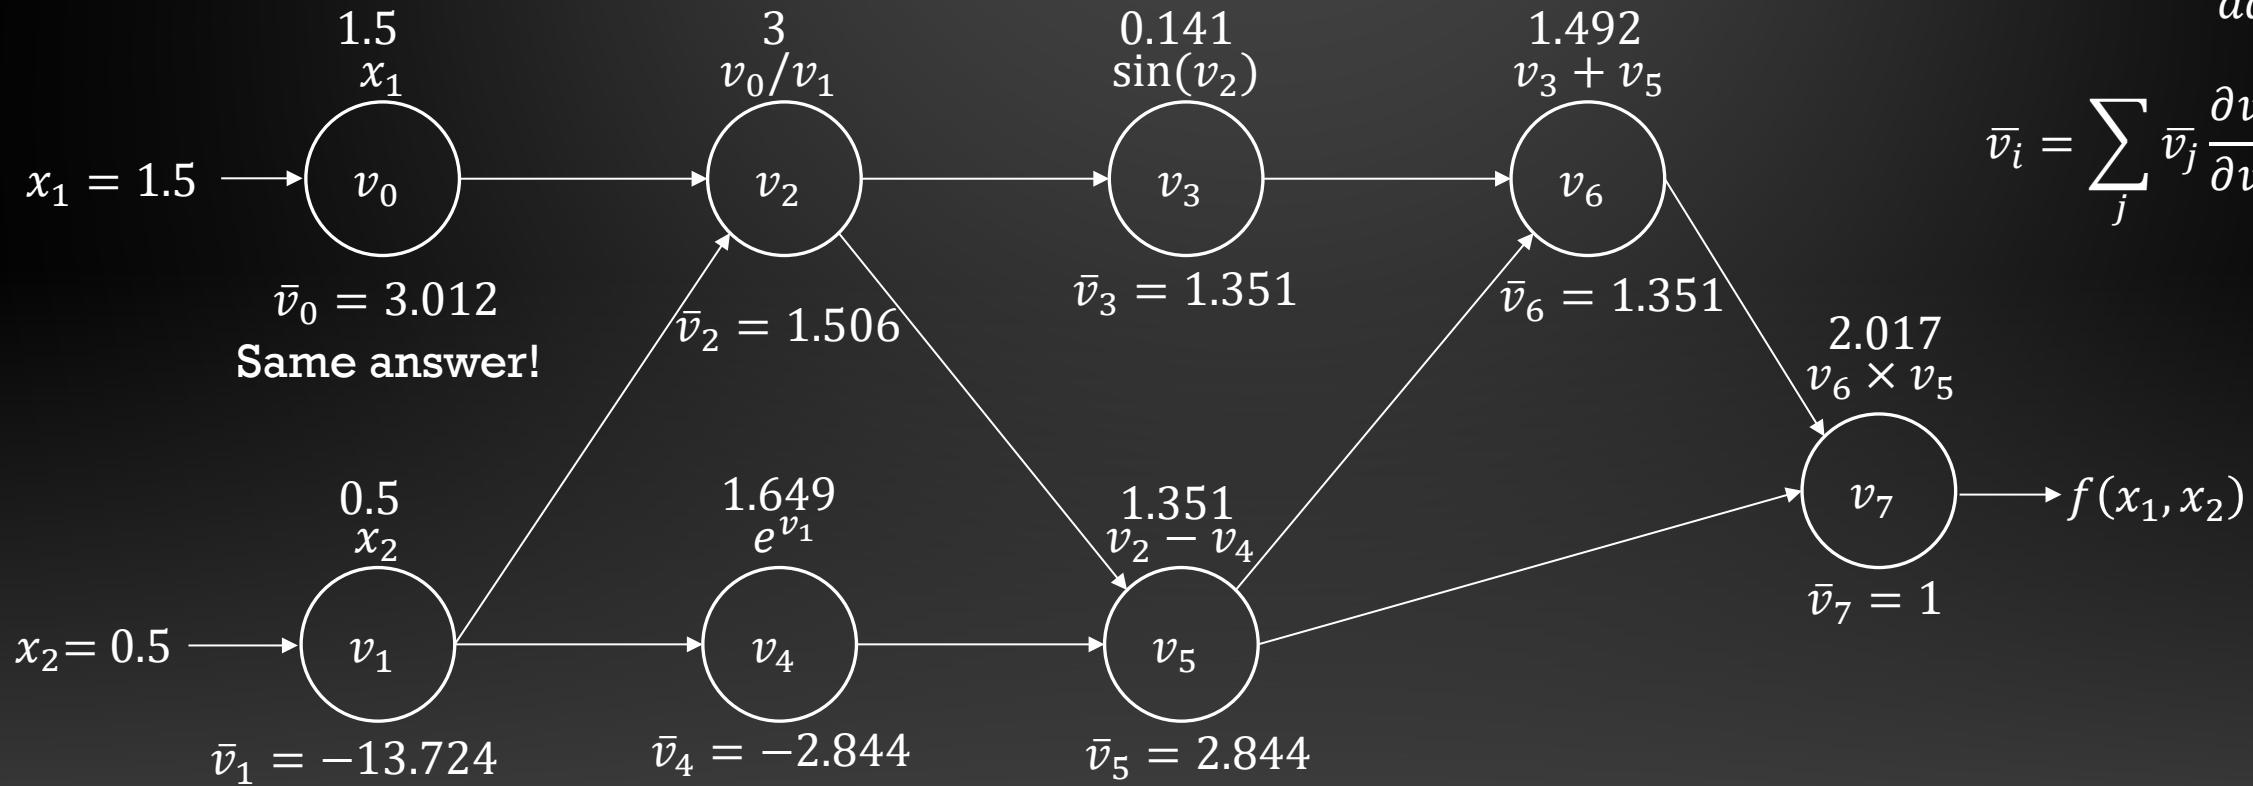
*function value*



*adjoint (v̄\_i)*

$$\bar{v}_i = \sum_j \bar{v}_j \frac{\partial v_j}{\partial v_i} \quad \text{for all } j \text{ ahead of } i$$

$$\bar{v}_0 = \bar{v}_2 \frac{\partial v_2}{\partial v_1} = 1.506 * \frac{1}{0.5} = 3.012$$



# REVERSE MODE AUTO-DIFF

Easily find partials wrt. all parameters for each output

Works well for ML/RL where  $N$  parameters  $\gg N$  outputs

Can be memory intensive as all intermediate states must be stored

Computationally relatively cheap (maximum 6x the cost of forward pass)

# SUMMARY

Neural Networks effective for approximating Q values

However, adjustments must be made for improved stability

Automatic differentiation is a useful tool for finding partial derivatives of programs

Reverse mode auto-diff more suited to ML/RL applications